

CS 202

Advanced Operating Systems

Winter 26

Lecture 3: OS Architecture

Instructor: Chengyu Song

What is an OS?


- Operating System
 - ◆ (of a person) control the functioning of (a machine, process, or system)
"a shortage of workers to operate new machines"
 - ◆ a set of things working together as parts of a mechanism or an interconnecting network
"the state railroad system"
 - ◆ a set of principles or procedures according to which something is done
"a multiparty system of government"
- The software that supports a computer's basic functions, such as scheduling tasks, executing applications, and controlling peripherals.

Topics for today

- ❑ What is an operating system?
- ❑ What is OS architecture / organization?
- ❑ Why does OS architecture / organization matter?
- ❑ Typical OS architectures


What is OS architecture?

- The design and organization of the software layer that defines the interface between applications and physical resources



OS architecture

- How OS functionalities are organized, and how components interact
- An earlier example:



Topics for today

- ❑ What is an operating system?
- ❑ What is OS architecture / organization?
- ❑ Why does OS architecture / organization matter?
- ❑ Typical OS architectures

Fundamental Issues

- The fundamental issues in OS designs are:
 - ◆ **Functionality**: how to enable the users/programs to do more?
 - ◆ **Performance**: how to do better?
 - ◆ **Management**: how to allocate and schedule resources for tasks?
 - ◆ **Protections**: how to make sure things won't go wrong?
 - ◆ **Communication**: how to enable collaboration?
 - ◆ **Security**: how to create a safe environment?
 - ◆ **Reliability and fault tolerance**: how to mask failures?

Why does OS architecture matter?

- Flexibility / Extensibility
 - ◆ How to support different hardware?
 - ◆ How to enable new hardware to applications?
 - ◆ How to adapt to different needs of applications?
- Performance
 - ◆ Does the structure facilitate good performance?
 - ◆ Scalability (to more applications / hardware resources)
 - ◆ Responsiveness (latency of event responses)
- Protection
 - ◆ User (mistakes, attacks)
 - ◆ System (bugs, faults, attacks)

U/U	U/S
S/U	S/S

Extensibility

- How to support different hardware?
 - ◆ CPU / ISA
 - ◆ MMU / memory access
 - ◆ Disk
 - ◆ Devices
- How to enable new hardware to applications?
 - ◆ Known hardware types / functionalities
 - ◆ New hardware types / functionalities

Extensibility

- How to adapt to different needs of applications?
 - ◆ Compile-time customization
 - » Linux kconfig
 - ◆ Boot-time configuration
 - » Boot options
 - ◆ Runtime reconfiguration
 - » /sys, /proc/sys, prctl, personality, priority
 - ◆ Programmable kernel extensions
 - » bpf, ebpf

One size does not fit all!

Performance

- What is bad for performance?
 - ◆ Copy (data transfer)
 - ◆ Wait (dependencies)
 - ◆ Cache misses
 - ◆ Checks
- Examples
 - ◆ Syscall: CPU context switch (saving / loading), copying args, checking args
 - ◆ Process switch: address space switch (TLB / predictors)
 - ◆ Swap / thrashing

Protection

- How to protect users and system?
 - ◆ Isolation (privilege mode, address spaces)
 - ◆ Verification (eliminate bugs / mistakes)
 - ◆ Checks (access control, defensive programming)

Fundamental Issues


operation

- The fundamental issues in ~~OS~~ designs are:

- ◆ **Functionality**: how to enable the users/programs to do more? → **add more code**
- ◆ **Performance**: how to be faster? → **remove the boundaries, eliminate checks**
- ◆ **Management**: how to allocate and schedule resources for tasks?
- ◆ **Protections**: how to make sure things won't go wrong? → **simple, isolations, checks**
- ◆ **Communication**: how to enable collaboration?
- ◆ **Security**: how to create a safe environment?
- ◆ **Reliability and fault tolerance**: how to mask failures?

Fundamental Trade-offs

- How to serve more customers?
 - ◆ Offer more products / services
 - ◆ Hire more people
- How to teach more students?
 - ◆ Offer more degrees
 - ◆ Establish more departments
 - ◆ Hire more people
- How to serve more citizens?
 - ◆ Offer more services
 - ◆ Establish more departments
 - ◆ Hire more people



- **Bureaucracy**
- **Bloat**
- **Inefficiency**
- **Slow to adapt**
- **Corruptions**


Topics for today

- What is an operating system?
- What is OS architecture / organization?
- Why does OS architecture / organization matter?
- Typical OS architectures

Typical OS architectures


Monolithic

- UNIX, Linux, Windows




Micro-kernel

- L4, Mach, QNX




LibraryOS

- exokernel, graphene





Typical OS architectures


- Multi-kernels
 - ◆ Barrelyfish, LegoOS




- Hypervisors
 - ◆ VMware ESX, Xen, KVM




AMAZON




GOOGLE




FACEBOOK




MICROSOFT




APPLE



ORACLE



Trade-offs of OS architectures



Trade-offs of OS architectures

- Monolithic kernels
 - ◆ Good performance
 - ◆ Okay protections (has isolation, too big to be bug free, hard to verify)
 - ◆ Bad extensibility
- Microkernels
 - ◆ Very good protection (isolation, small, verifiable)
 - ◆ Okay extensibility
 - ◆ Bad performance
- LibraryOS
 - ◆ Very good extensibility
 - ◆ Good performance
 - ◆ Bad protection (no isolation between U/libOS, exposes everything)


Exokernel

- Motivations: fixed high-level abstractions are hurting applications
 - ◆ Fixed high-level abstractions hurt application performance by denying domain-specific optimizations
 - » Scheduler, cache eviction policy
 - » Evidences: Cao et al. [1 O] measured that application-controlled file caching can reduce application running time by as much as 45%.
 - ◆ Fixed high-level abstractions discourages changes to existing abstractions and limit the functionality of applications
 - » Syscalls rarely change, if ever
 - » Evidences: few research ideas have been adopted
 - ◆ Fixed high-level abstractions hide information from applications and prevent it from making good decisions
 - » Low-level exceptions, interrupts, raw device I/O
 - » Evidences: database implementors must struggle to emulate random-access record storageon top of tile systems [47].

Exokernel

- Core idea: **Separating Protection from Management**
 - ◆ Exokernel: provides protection
 - ◆ LibraryOS: implements management and high-level abstractions
- Solution: **securely** multiplexes available **hardware resources** through a **low-level** interface
 - ◆ Low-level: even more flexible / aggressive than microkernel
 - ◆ No emulation: direct exposing hardware (including naming): much lower overhead than virtual machine monitors
 - ◆ Securely: trust but verify

Exokernel: architecture



Exokernel: Design Principle

- **Securely expose hardware**
 - ◆ An exokernel strives to safely export all privileged instructions, hardware DMA capabilities, and machine resources
- **Expose allocation**
 - ◆ An exokernel should allow library operating systems to request specific physical resources
- **Expose names**
 - ◆ An exokernel should export physical names
- **Expose Revocation**
 - ◆ An exokernel should utilize a visible resource revocation protocol so that a well-behaved libraryOS can perform effective application-level resource management
 - ◆ All revocations, including CPU (i.e., explicit context switch)
- **Arbitration:** same as other OS

Exokernel: secure bindings

- Tracks resource ownership
 - ◆ Allocation to libraryOS at *bind time*
 - ◆ Free to use afterwards, no more authorization
- Protection (ownership) checks are simple operations performed by the kernel
 - ◆ Hardware mechanisms (if available, e.g., capability tokens)
 - » Memory: TLB, page tables, DMA
 - ◆ Software caching (bookkeeping)
 - ◆ Downloadable application code
 - » Packet filter
- Allows protection without understanding

Exokernel: secure binding example (1)

- Virtual-to-Physical address translation
 - ◆ Protecting which physical pages can be accessed by a libraryOS
- TLB-based
 - ◆ LibraryOS decide the virt-to-phy mapping
 - ◆ LibraryOS decide the TLB replacement policy
 - ◆ Binding presented to exokernel
 - ◆ Exokernel checks ownership, if passed, installs/removes entries in hardware TLB
 - ◆ LibraryOS can access physical memory freely
- Page-table based
 - ◆ More challenging, Xen, hardware virtualization

Exokernel: secure binding example (2)

- Packet filtering
 - ◆ Handle packet distributing (who handles?) and processing (e.g., firewall)
 - ◆ In-kernel processing is much faster than polling / querying
- Manage packet handling from library (app policy)
 - ◆ Packet filter code provided by library (app)
 - ◆ In the form of downloadable code
 - ◆ Binding (code) presented to exokernel
 - ◆ Exokernel verifies the code
 - ◆ Exokernel installs app's packet filter code (e.g., to hardware NIC)

Exokernel: hypotheses

- The design and implementation of exokernel is built upon four hypotheses that must be validated through empirical evaluation
 - ◆ Exokernels can be very efficient
 - ◆ Low-level, secure multiplexing of hardware resources can be implemented efficiently
 - ◆ Traditional operating system abstractions can be implemented efficiently at application level
 - ◆ Applications can create special-purpose implementations of these abstractions
- Q: how to design experiments?