

Lab1.3: Copy-on-Write fork - Coding

Lab1.3: Copy-on-Write fork - Coding

Your task is to implement copy-on-write fork in the xv6 kernel. You are done if your modified kernel executes both the `cowtest` and `usertests` programs successfully.

Link to the GitHub Classroom assignment: <https://classroom.github.com/a/UDtW9Z8r>.

To start, clone the private lab assignment repo from GitHub classroom. Refer to the [setup tutorial page](#) for help.

```
$ git clone git@github.com:UCR-CS202/lab-...
```

Virtual memory provides a level of indirection: the kernel can intercept memory references by marking PTEs invalid or read-only, leading to page faults, and can change what addresses mean by modifying PTEs. There is a saying in computer systems that any systems problem can be solved with a level of indirection. This lab explores an example: copy-on-write fork.

The Problem

The `fork()` system call in xv6 copies all of the parent process's user-space memory into the child. If the parent is large, copying can take a long time. Worse, the work is often largely wasted: `fork()` is commonly followed by `exec()` in the child, which discards the copied memory, usually without using most of it. On the other hand, if both parent and child use a copied page, and one or both writes it, the copy is truly needed.

The Solution

Your goal in implementing copy-on-write (COW) `fork()` is to defer allocating and copying physical memory pages until the copies are actually needed, if ever. COW `fork()` creates just a pagetable for the child, with PTEs for user memory pointing to the parent's physical pages. COW `fork()` marks all the user PTEs in both parent and child as read-only. When either process tries to write one of these COW pages, the CPU will force a page fault.

The kernel page-fault handler detects this case, allocates a page of physical memory for the faulting process, copies the original page into the new page, and modifies the relevant PTE in the faulting process to refer to the new page, this time with the PTE marked writeable. When the page fault handler returns, the user process will be able to write its copy of the page.

COW `fork()` makes freeing of the physical pages that implement user memory a little trickier. A given physical page may be referred to by multiple processes' page tables, and should be freed only when the last reference disappears. In a simple kernel like xv6 this bookkeeping is reasonably straightforward, but in production kernels this can be difficult to get right; see, for example, “Patching until the COWs come home.”

Implement copy-on-write fork

Your task is to implement copy-on-write `fork` in the xv6 kernel. You are done if your modified kernel executes both the `cowtest` and `usertests -q` programs successfully.

To help you test your implementation, we've provided an xv6 program called `cowtest` (source in `user/cowtest.c`). `cowtest` runs various tests, but even the first will fail on unmodified xv6. Thus, initially, you will see:

```
$ cowtest
simple: fork() failed
$
```

The “simple” test allocates more than half of available physical memory, and then `fork()`s. The fork fails because there is not enough free physical memory to give the child a complete copy of the parent's memory.

When you are done, your kernel should pass all the tests in both `cowtest` and `usertests -q`. **If you haven't thought about usertests in your evaluation plan, you missed.** That is:

```
$ cowtest
simple: ok
simple: ok
three: ok
three: ok
three: ok
file: ok
forkfork: ok
ALL COW TESTS PASSED
$ usertests -q
...
ALL TESTS PASSED
$
```

Here's a reasonable plan of attack, **check your previous design against it.**

1. Modify `uvncpy()` to map the parent's physical pages into the child, instead of allocating new pages. Clear `PTE_W` in the PTEs of both child and parent for pages that have `PTE_W` set.
2. Modify `vmfault()` to recognize page faults. When a write page-fault occurs on a COW page that was originally writeable, allocate a new page with `kalloc()`, copy the old page to the new page, and install the new page in the PTE with `PTE_W` set. Pages that were originally read-only (not mapped `PTE_W`, like pages in the text segment) should remain read-only and shared between parent and child; a process that tries to write such a page should be killed.
3. Ensure that each physical page is freed when the last PTE reference to it goes away -- but not before. A good way to do this is to keep, for each physical page, a "reference count" of the number of user page tables that refer to that page. Set a page's reference count to one when `kalloc()` allocates it. Increment a page's reference count when fork causes a child to share the page, and decrement a page's count each time any process drops the page from its page table. `kfree()` should only place a page back on the free list if its reference count is zero. It's OK to keep these counts in a fixed-size array of integers. You'll have to work out a scheme for how to index the array and how to choose its size. For example, you could index the array with the page's physical address divided by 4096, and give the array a number of elements equal to highest physical address of any page placed on the free list by `kinit()` in `kalloc.c`. Feel free to modify `kalloc.c`(e.g., `kalloc()` and `kfree()`) to maintain the reference counts.
4. Modify `copyout()` to use the same scheme as page faults when it encounters a COW page.

Some hints:

- It may be useful to have a way to record, for each PTE, whether it is a COW mapping. You can use the RSW (reserved for software) bits in the RISC-V PTE for this.
- `usertests -q` explores scenarios that `cotest` does not test, so don't forget to check that all tests pass for both.
- Some helpful macros and definitions for page table flags are at the end of `kernel/riscv.h`.
- If a COW page fault occurs and there's no free memory, the process should be killed.

Submissions

As mentioned in Lab1.2, you need to submit two sets of code:

1. A test program `cow_perf` that outputs performance metrics to show COW indeed improved performance.
2. The actual COW implementation in the xv6 kernel that can pass all the tests listed above.

Create two commits. First commit your performance evaluation code and any kernel changes made to collect performance metrics. We will collect the performance metrics before COW is added first. Then commit your COW implementation.

`cow_perf` requirements (autograded):

- `cow_perf` must be runnable from the xv6 shell (i.e., included in `UPROGS`), and should print a single non-negative integer (digits only) on its own line.
- Pick a metric that should change with COW. You choose the direction by editing `cow_perf_direction`:
 - `cow_perf_direction` must contain either `less` or `greater`.
 - The autograder checks that the number at your final COW commit (`HEAD`) is strictly less/greater than the number at your chosen baseline commit.
- You choose the baseline commit by editing `cow_baseline_commit`:
 - `cow_baseline_commit` must contain a git commit hash (full or abbreviated) that exists in your repo.
 - Typically, set this to your first commit (your `cow_perf` commit) before adding COW.
- Your `cow_perf` should run on both the baseline commit and your final COW commit (`HEAD`). In particular, you should not add new syscalls or kernel-only instrumentation between the baseline commit and `HEAD`.

Example of a `user/cow_perf.c` (prints a constant, so baseline and `HEAD` will match and the improvement check will fail):

```
#include "kernel/types.h"
#include "user/user.h"
int main(void) { printf("123\n"); exit(0); }
```

Submit the lab

Submit

We will use GitHub Classroom for submissions.

Please merge the GitHub Classroom PR that updates the grader scripts (it includes the `cow_perf` autograder checks), then pull the updated files locally before running `make grade`. Deadline update: Lab 1 submission deadline is extended to Feb 8 (Sunday).

Commit your changes and push them to your Classroom repository. A GitHub Action will run the tests and report whether your implementation is correct.

- Please run `make grade` locally before pushing.
- The autograder runs automatically in GitHub Actions and syncs your score to the corresponding Canvas assignment.
- If the Action fails, open the failed workflow run and check its output to see which tests you're missing and why

you didn't receive full credit.

- To enable Canvas sync, put your UCR email in `email.txt` (single line like `netid@ucr.edu`). If your Canvas submission can't be found, the CI log will ask you to contact the TA.

Reference

<https://pdos.csail.mit.edu/6.1810/2025/labs/cow.html>