

# **CS 202**


# **Advanced Operating Systems**

**Winter 26**

Lecture 2: Historical Perspective

Instructor: Chengyu Song


# Phase 1: batch systems (1955-1970)



- Computers expensive; people cheap
  - ◆ Use computers efficiently – move people away from machine
- OS in this period became a program loader
  - ◆ Loads a job, runs it, outputs result, then moves on to next
  - ◆ More efficient use of hardware but increasingly difficult to debug
    - » Still less bugs ☺

# Phase 1, problems

- Utilization is low (one job at a time)
- No protection between jobs
  - ◆ But one job at a time, so what can go wrong?
- Scheduling
- Coordinating concurrent activities
- People time is still being wasted
- Operating Systems didn't really work
  - ◆ The mythical man month
  - ◆ Birth of software engineering



# Phase 2: 1970s

- Computers and people are expensive
  - ◆ Help people be more productive
- **Interactive time-sharing:** let many people use the same machine at the same time
- Emergence of minicomputers
  - ◆ Terminals are cheap
- Persistence: keep data online on fancy file systems

# Modern architecture support for OS

- Manipulating privileged machine state
  - ◆ Protected instructions
  - ◆ Manipulate device registers, TLB entries, etc.
  - ◆ Controlling access
  - ◆ Memory isolation and translation
- Generating and handling “events”
  - ◆ Interrupts, exceptions, system calls, etc.
  - ◆ Respond to external events
  - ◆ CPU requires software intervention to handle fault or trap
- Other stuff
  - ◆ Mechanisms to handle concurrency, isolation, virtualization ...

# Protected Instructions

- OS must have exclusive access to hardware and critical data structures
- Only the operating system can
  - ◆ Directly access I/O devices (disks, printers, etc.)
    - » Security, fairness (why?)
  - ◆ Manipulate memory management state
    - » Page table pointers, page protection, TLB management, etc.
  - ◆ Manipulate protected control registers
    - » Kernel mode, interrupt level
  - ◆ Halt instruction (why?)

# Privilege Mode

- Hardware restricts privileged instructions to OS
- Q: How does the HW know if the executed program is OS?
  - ◆ HW must support (at least) two execution modes: OS (kernel) mode and user mode
- Mode kept in a status bit in a protected control register
  - ◆ User programs execute in user mode
  - ◆ OS executes in kernel mode (OS == “kernel”)
  - ◆ CPU checks mode bit when protected instruction executes
  - ◆ Attempts to execute in user mode trap to OS


# Switching back and forth

- Going from higher privilege to lower privilege
  - ◆ Easy: can directly modify the mode register to drop privilege
- But how do we escalate privilege?
  - ◆ Special instructions to change mode
    - » System calls (int 0x80, syscall, svc)
    - » Saves context and invokes designated handler
      - You jump to the privileged code; you cannot execute your own
      - » OS checks your syscall request and honors it only if safe
    - ◆ Or, some kind of event happens in the system


# Memory Isolation

- OS must be able to protect programs from each other
- OS must protect itself from user programs
- OS may or may not protect user programs from itself
- Memory management unit (MMU)
  - ◆ Hardware unit provides memory protection mechanisms
  - ◆ Virtual memory
  - ◆ Segmentation
- Manipulating memory management hardware uses protected (privileged) operations

# Example memory protection



# Core i7 Level 1-3 Page Table Entries



**P:** Child page table present in physical memory (1) or not (0).

**A:** Reference bit (set by MMU on reads and writes, cleared by software).

**R/W:** Read-only or read-write access permission for all reachable pages.

**PS:** Page size either 4 KB or 2 MB (defined for Level 1 PTEs only).

**U/S:** user or supervisor (kernel) mode access permission for all reachable pages.

**G:** Global page (don't evict from TLB on task switch)

**WT:** Write-through or write-back cache policy for the child page table.

**Page table physical base address:** 40 most significant bits of physical page table address (forces page tables to be 4KB aligned)


**CD:** Caching disabled or enabled for the child page table.

**XD:** Non-executable pages

# Types of Arch Support

- Manipulating privileged machine state
  - ◆ Protected instructions
  - ◆ Manipulate device registers, TLB entries, etc.
  - ◆ Controlling access
  - ◆ Memory isolation and translation
- Generating and handling “events”
  - ◆ Interrupts, exceptions, system calls, etc.
  - ◆ Respond to external events
  - ◆ CPU requires software intervention to handle fault or trap

# Review: Computer Organization




# Events

- An event is an “unnatural” change in control flow
  - ◆ Events immediately stop current execution
  - ◆ Changes mode, context (machine state), or both
- Most events are handled by the OS kernel, which defines a handler for each event type
  - ◆ Event handlers always execute in kernel mode
  - ◆ The specific types of events are defined by the machine
- Once the system is booted, OS is one big event handler
  - ◆ all entry to the kernel occurs as the result of an event

# Context Switch


- An event is an “unnatural” change in execution, so there are a few important questions
  - ◆ How to make sure the event can be properly handled?
    - » Who (what code) should handle the current event
    - » What preparation is necessary to handle the current event
  - ◆ Once the event is done handling, how do we go back to previous execution?
  - ◆ Essentially, how to switch between different tasks?
- **Context Switch** is the key concept in understanding this, when event happens, the hardware will
  - ◆ Save the current execution context (CPU registers)
  - ◆ Setup the context for the handler (entry point of code, stack, privilege mode, etc.)
  - ◆ (after done) Restore the previous context

# Handling events – Interrupt Vector Table



# Categorizing Events

- Two *kinds* of events: **synchronous** and **asynchronous**
- Sync events are caused by executing instructions
  - ◆ Example?
- Async events are caused by an external event
  - ◆ Example?



# Categorizing Events

- This gives us a convenient table:

	Unexpected	Scheduled
Synchronous	fault/exception	syscall
Asynchronous	interrupt	signal (and timer)

- Terms may be slightly different by OS and architecture
  - E.g., POSIX signals, async system traps, async or deferred procedure calls

# Signals

- Software interrupts, from
  - ◆ Users (terminals, keyboard)
  - ◆ Other processes
  - ◆ OS (process' own execution)

Signal	ID	Description	Default Action
SIGINT	2	<b>Interrupt.</b> Sent when you press Ctrl+C. It politely tells the process to stop.	Terminate
SIGTERM	15	<b>Terminate.</b> The default signal sent by the kill command. It asks the process to clean up and exit.	Terminate
SIGKILL	9	<b>Kill.</b> The "nuclear option." It effectively pulls the power cord on the process. It cannot be ignored or handled.	Force Kill
SIGSEGV	11	<b>Segmentation Fault.</b> The process tried to read/write memory it doesn't own.	Dump Core
SIGHUP	1	<b>Hangup.</b> Sent when a terminal closes. Daemons often use this as a trigger to <b>reload configuration files</b> without stopping.	Terminate
SIGSTOP	19	<b>Stop.</b> Pauses the process (like pressing pause on a movie). Cannot be ignored.	Pause
SIGCONT	18	<b>Continue.</b> Resumes a process previously paused by SIGSTOP.	Resume

# Other Staff

- Support for synchronization
  - ◆ How to build locks?
    - » Atomic operations
    - » Disabling interrupts
  - ◆ How to support memory consistency
    - » One CPU writes, visible to other CPUs
    - » How to keep caches consistent
  - ◆ How to communicate between CPU cores?

# The TENEX System (Hardware/OS)

- Design Philosophy
  - ◆ The primary goal was to create a virtual machine that exceeded the physical limitations of the processor (PDP-10)
- The BBN Pager
  - ◆ Innovative hardware that provided individual mapping for 512-word pages → [MMU](#)
- Associative Registers
  - ◆ The pager used these to store mapping information, allowing for rapid address translation → [TLB](#)
- Large Address Space:
  - ◆ It provided a paged virtual address space [equal to or greater](#) than the processor's addressing capability with full protection and sharing

# TENEX: Memory Management

- Copy-on-Write (COW)
  - ◆ TENEX introduced the "copy-on-write" facility, allowing processes to share large address spaces and only create private copies when a page is actually modified
- The Working Set Principle
  - ◆ Memory management was based on Peter Denning's "working set principle," aiming to identify the "balance set" of processes that could fit in core memory simultaneously to minimize thrashing
- Core Status Table
  - ◆ The hardware maintained a record of page activity (references and writes) in a core status table, which the software used to make informed page replacement decisions
- Demand Paging
  - ◆ No pages were preloaded; the system relied entirely on demand paging, making the paging traps invisible to the user process

# TENEX: File System

- File-to-Memory Mapping
  - ◆ TENEX allowed a virtual memory slot to contain a pointer to a page from a file, effectively integrating the file system into the virtual address space → [mmap](#)
- Symbolic Naming and Versioning
  - ◆ It featured a multilevel symbolic directory structure and automatic version numbering for files. Every time a file was written, the system incremented the version number to prevent accidental data loss
- Uniform I/O
  - ◆ The system treated all external devices and data streams through a uniform interface, allowing I/O to be handled consistently regardless of the device → [UNIX](#)
- Thawed vs. Unthawed Access
  - ◆ The system managed file sharing through "thawed" access (allowing multiple writers/readers without consistency guarantees) or "unthawed" access (enforcing single-writer/multiple-reader consistency)

# TENEX: User Interactions

- User-Centric Design
  - ◆ A major goal was "good human engineering," focusing on the TENEX Executive (EXEC) command interpreter → [shell and coreutils](#)
- Command Completion
  - ◆ The EXEC introduced the use of the ESC key for command and filename completion, as well as the "?" key for context-sensitive help
- Full-Duplex Interaction
  - ◆ The system was optimized for full-duplex terminals, allowing for "intimate interaction" where users could type input while the machine was still outputting data
- Pseudo-Interrupt System
  - ◆ This allowed processes to receive asynchronous signals from other processes or terminals, facilitating sophisticated terminal control → [signals](#)

# Unix appears

- Ken Thompson, who worked on MULTICS, wanted to use an old PDP-7 laying around in Bell labs
- He and Dennis Richie built a system designed by programmers for programmers
- Originally in assembly. Rewritten in C
  - ◆ In their paper describing UNIX, they defend this decision!
  - ◆ However, this is a new and important advance: portable operating systems!
- Shared code with everyone (particularly universities)