

CS 202

Advanced Operating Systems

Winter 26

Lecture 1: Course Introduction

Instructor: Chengyu Song

Associate Professor


Department of Computer Science and Engineering

Teaching Staff



Chengyu Song

I am an Associate Professor in CSE
Office hours TTh 11:00pm-12:00pm or by appointment
• In person and online



TA for Labs

Jiajun Chen (PhD student)
Office hours and links on Canvas
Leads for Labs

Expectations

- Advanced
 - ◆ You should have already learnt about basic operating system concepts
 - » We will do a quick review
 - ◆ **Read papers** (together)
 - » Advanced topics
 - ◆ **Discussion and active class participation** (10% of final grade)
 - ◆ More challenging labs

What is this course about?

- How has the role of the operating system **evolved** over time?
 - ◆ How does the past inform the present?
- What are the **principles** that underlie Operating Systems?
- What are current and future **trends** in OS?
- Make it real: lab assignments to get some experience with OS development
- *Get you ready to do systems research*

Some topics we will cover

- Operating Systems models and how they evolved
 - ◆ Extensibility, protection, performance
- Concurrency:
 - ◆ Synchronization and Scheduling
 - ◆ Multicore OS
- File systems:
 - ◆ Local, networked, distributed, internet scale
- Other advanced topics

Class format

- For every topic:
 - ◆ Some overview
 - ◆ **Read related book chapters** before the class
 - ◆ **Discuss research papers**
- Research papers:
 - ◆ Critique for some required papers (1-2 paper per topic)
 - ◆ You are responsible for required papers and material discussed in class

Reading Research Papers

- Guidelines for reading papers
 - ◆ Make sure to identify authors' goals and assumptions. Not always directly stated.
 - ◆ Look for high-level takeaways.
 - ◆ **Simulate the whole process in your head**
 - ◆ Follow a multi-pass reading strategy
 - » Pass1: Get overview. Pass2: Read details and make notes. Pass3: Re-read details to evaluate.
 - ◆ Think how techniques used are applicable today. Identify extensions.

Questions while Reading Papers

- What are the primary goals (hypothesis)?
 - ◆ 2 sentence elevator pitch
- Why did the authors do what they did the way they did?
 - ◆ Understand motivation for design
- What were the driving forces for the paper at the time it was written?
- What parts still hold? What parts don't?
- How did they experiment to support their ideas?
- **Write the review by yourself**

Projects

- 3 Lab Assignments
 - ◆ Modifications on xv6 (<https://github.com/mit-pdos/xv6-riscv>)
 - ◆ Related to the topics discussed
- All labs are done individually
- Submission
 - ◆ Design documents & test plans
 - ◆ Code & tests
 - ◆ Short demo
- Deadlines are **hard!**
- Late policy
 - ◆ 10% penalty per day; up to 60% late penalties

Course Materials

- Textbook: Apraci-Dessau and Apraci-Dessau, **OS, 3 easy pieces**
 - ◆ <https://pages.cs.wisc.edu/~remzi/OSTEP/>
 - ◆ Please read, most chapters are only a few pages
- Lecture notes and Papers: on Canvas
- Other good books:
 - ◆ Anderson and Dahlin, ***Operating Systems: Principles and Practice (recommended)***
 - ◆ Silberschatz, Galvin, and Gagne, ***Operating System Concepts***, John Wiley and Sons, 8th Edition **(recommended)**

Grading Policy

- Lab Assignments: 40%
 - ◆ Lab 1: 10%, Lab 2: 15%, Lab 3:15%
- Reading and critiquing papers: 20%
- Attendance & Class Participation: 10%
- Final (cumulative, closed book): 30%

Cheating and Plagiarism

- Your code and critiques may be compared against each other
- Academic integrity policies and procedures
 - ◆ <https://conduct.ucr.edu/policies/academic-integrity-policies-and-procedures>
- What consequences can happen?
 - ◆ Will be reported to the Academic Integrity Committee (AIC)
 - ◆ Grade penalty
 - ◆ Course failure
 - ◆ Academic sanctions: probation, suspension, dismissal

Policy on AI

- AI is replacing many entry-level positions
 - ◆ You do not want to become replaceable
 - » If you just copy&paste whatever AI told you, you're replaceable, elementary school kids can do this
 - » If you stop thinking, you become replaceable, as we're teaching AI to think/reason
- Learning to work with AI is necessary
 - ◆ Use AI to help you understand research papers better, but write the review by yourself and using your own words
 - ◆ Use AI to help you understand existing code better, but put your own thoughts into design and testing
 - ◆ Use AI to help you code, but follow your design and know how to test and reason

Quick Review of Operating System

Questions for today

- What is “operating”?
- What is an operating system?
- Why do we need operating systems?
- What do operating systems do?

Operating


- Ask ChatGPT: what is operating
- Get things done (hide the details)
- Management of resources
- Optimization (reduce cost, improve deliveries, robustness)

Operating System

- Why we have computers?
- Operating system: the COO of computer systems
 - ◆ Get things done (hide the details)
 - ◆ Management of resources
 - ◆ Optimization

Goals of an OS/Platform

- **Enabling:** allow users and developers to use (new) hardware to do new things that are **not possible before**



- **Usability:** make it **easy** to use / **program with**




- **Efficiency:** are we **wasting** our money/time/energy?

- **Fairness and protection:** make sure good things will happen and bad things will not

Roles of an OS

- **Abstraction:** defines a set of logical resources (**objects**) and well-defined operations on them (**interfaces**)
 - ◆ **Why?** Easier app programming
 - ◆ Humans are good at abstraction instead of details
- **Virtualization:** Isolates and multiplexes physical resources via spatial and temporal sharing
 - ◆ **Why?** Easier programming, better hardware utilization
- **Control:**
 - ◆ **Why?** Fairness, performance, security, privacy, etc.

OS Abstractions



How to Share Computing/Processors

- Time sharing
 - ◆ Events and controlled/limited direct execution
 - ◆ Context switch (user <-> kernel, thread1 <-> thread2)
 - ◆ Scheduling (OS controls)
 - ◆ Synchronization (application controls)
- Parallel and concurrent
 - ◆ Parallel: (really/physically) at the same time, with none or minimal dependencies
 - ◆ Concurrent: (conceptually) at the same time, likely with shared resources
 - ◆ Consensus: who gets the lock / bitcoin

How to Share/Access Storages

- Translation
 - ◆ Virtual address space
 - » Virtual address → physical address
 - ◆ Disk and File Systems
 - » Name → file object (e.g., inode)
 - » Offset → disk blocks
- Cache
 - ◆ Program Locality
 - ◆ Memory Hierarchy

Translation methods

- Linear/contiguous mapping: $y = x + b$
 - ◆ Physical address = base register + virtual address
 - ◆ Disk block = starting block + offset / block size
 - ◆ Fast, easy to implement (no cache needed)
 - ◆ Inflexible, causes fragmentation (either internal or external)
- Indexed mapping: $y = \text{map}[x]$
 - ◆ Physical page number = page table[virtual page number]
 - ◆ File = directory[file name]
 - ◆ Disk block = inode[virtual block number]
 - ◆ IP address = DNS[domain name]
 - ◆ Could go through multiple levels
 - ◆ Flexible, easy to manage free space, no external fragmentation
 - ◆ Translation is slower (cache to the rescue)

Intel IA32-e Paging




Figure 4-8. Linear-Address Translation to a 4-KByte Page using IA-32e Paging


Source: Intel Architecture Software Developer Manuals

Path Name Translation

- Let's say you want to open “/one/two/three”
- What does the file system do?
 - ◆ Open directory “/” (well known, can always find)
 - ◆ Search for the entry “one”, get location of “one” (in dir entry)
 - ◆ Open directory “one”, search for “two”, get location of “two”
 - ◆ Open directory “two”, search for “three”, get location of “three”
 - ◆ Open file “three”


Unix Inodes

- Unix inodes implement an indexed structure for files
 - ◆ Also store metadata info (protection, timestamps, length, ref count...)
- Each inode contains 15 block pointers
 - ◆ First 12 are direct blocks (e.g., 4 KB blocks)
 - ◆ Then single, double, and triple indirect



Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently
- **Temporal locality:**
 - ◆ Recently referenced items are likely to be referenced again in the near future
- **Spatial locality:**
 - ◆ Items with nearby addresses tend to be referenced close together in time



Locality Example

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

- Data references
 - ◆ Reference array elements in succession (stride-1 reference pattern). Spatial locality
 - ◆ Reference variable **sum** each iteration. Temporal locality
 - Instruction references
 - ◆ Reference instructions in sequence.
 - ◆ Cycle through loop repeatedly. Spatial locality
 - ◆ Temporal locality

Cache

- | **Cache:** a smaller, faster storage that acts as a staging area for a subset of the data in a larger, slower storage.
 - ◆ The storage could be a software data structure (like index tables) or a hardware device (like spinning disk)
- | Why does cache work?
 - ◆ Because of **locality!**
 - » Hit fast storage much more frequently even though its smaller


General Cache Concepts

- Hit
 - ◆ Data needed is in the cache
- Miss
 - ◆ Data is not in the cache
 - ◆ Types of misses
 - » **Cold (compulsory) miss:** cache is empty
 - » **Conflict miss:** cache is not full but due to placement policy, the slot the data will be mapped to is occupied
 - » **Capacity miss:** the set of active cache blocks (**working set**) is larger than the cache size

Cache replacement policy

- Cache replacement policy: determine which data to remove when there is a miss
- Belady's algorithm
 - ◆ Idea: Replace the page that will not be used for the longest time in the future
 - ◆ Optimal but not practical: must predict the future
 - ◆ Serves as the baseline to measure other algorithms
- Common replacement algorithms
 - ◆ FIFO, Least Recently Used (LRU), access bit, LRU clock

An Example Memory Hierarchy



Examples of Caching in the Hierarchy

Cache Type	What is Cached?	Where is it Cached?	Latency (cycles)	Managed By
Registers	4-8 bytes words	CPU core	0	Compiler
TLB	Address translations	On-Chip TLB	0	Hardware
L1 cache	64-bytes block	On-Chip L1	1	Hardware
L2 cache	64-bytes block	On/Off-Chip L2	10	Hardware
Virtual Memory	4-KB page	Main memory	100	Hardware + OS
Buffer cache	Parts of files	Main memory	100	OS
Disk cache	Disk sectors	Disk controller	100,000	Disk firmware
Network buffer cache	Parts of files	Local disk	10,000,000	AFS/NFS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Security and Reliability

- Access control
 - ◆ DAC: access control list and capabilities
 - ◆ CS 165
- Reliability
 - ◆ Crash consistency: journaling file systems (**JFS**)
 - ◆ Disk failures: Redundant Array of Inexpensive Disks (**RAID**)