

# Big-data Systems: A Tour

Ahmed Eldawy

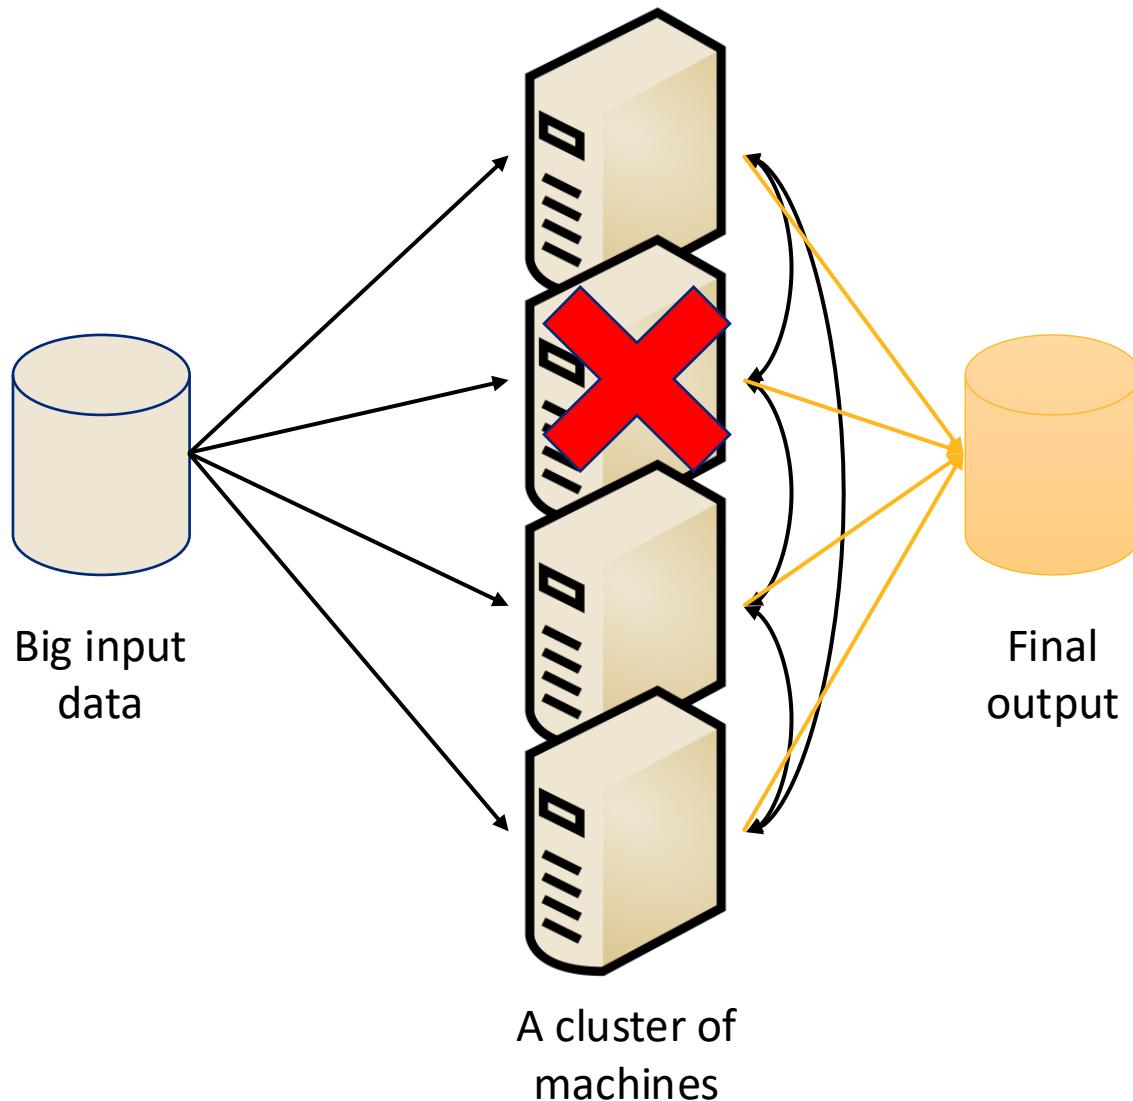
# Distributed Data Processing

- The idea of distributed databases is older than you might think

Richard Peebles, Eric G. Manning: A Computer Architecture for Large (Distributed) Data Bases. VLDB 1975: 405-427

- Distributed data structures and algorithms have always been around
- So, what is new?

# Distributed Data Processing



Data partitioning  
Load balancing  
Fault tolerance  
Synchronization

# MapReduce

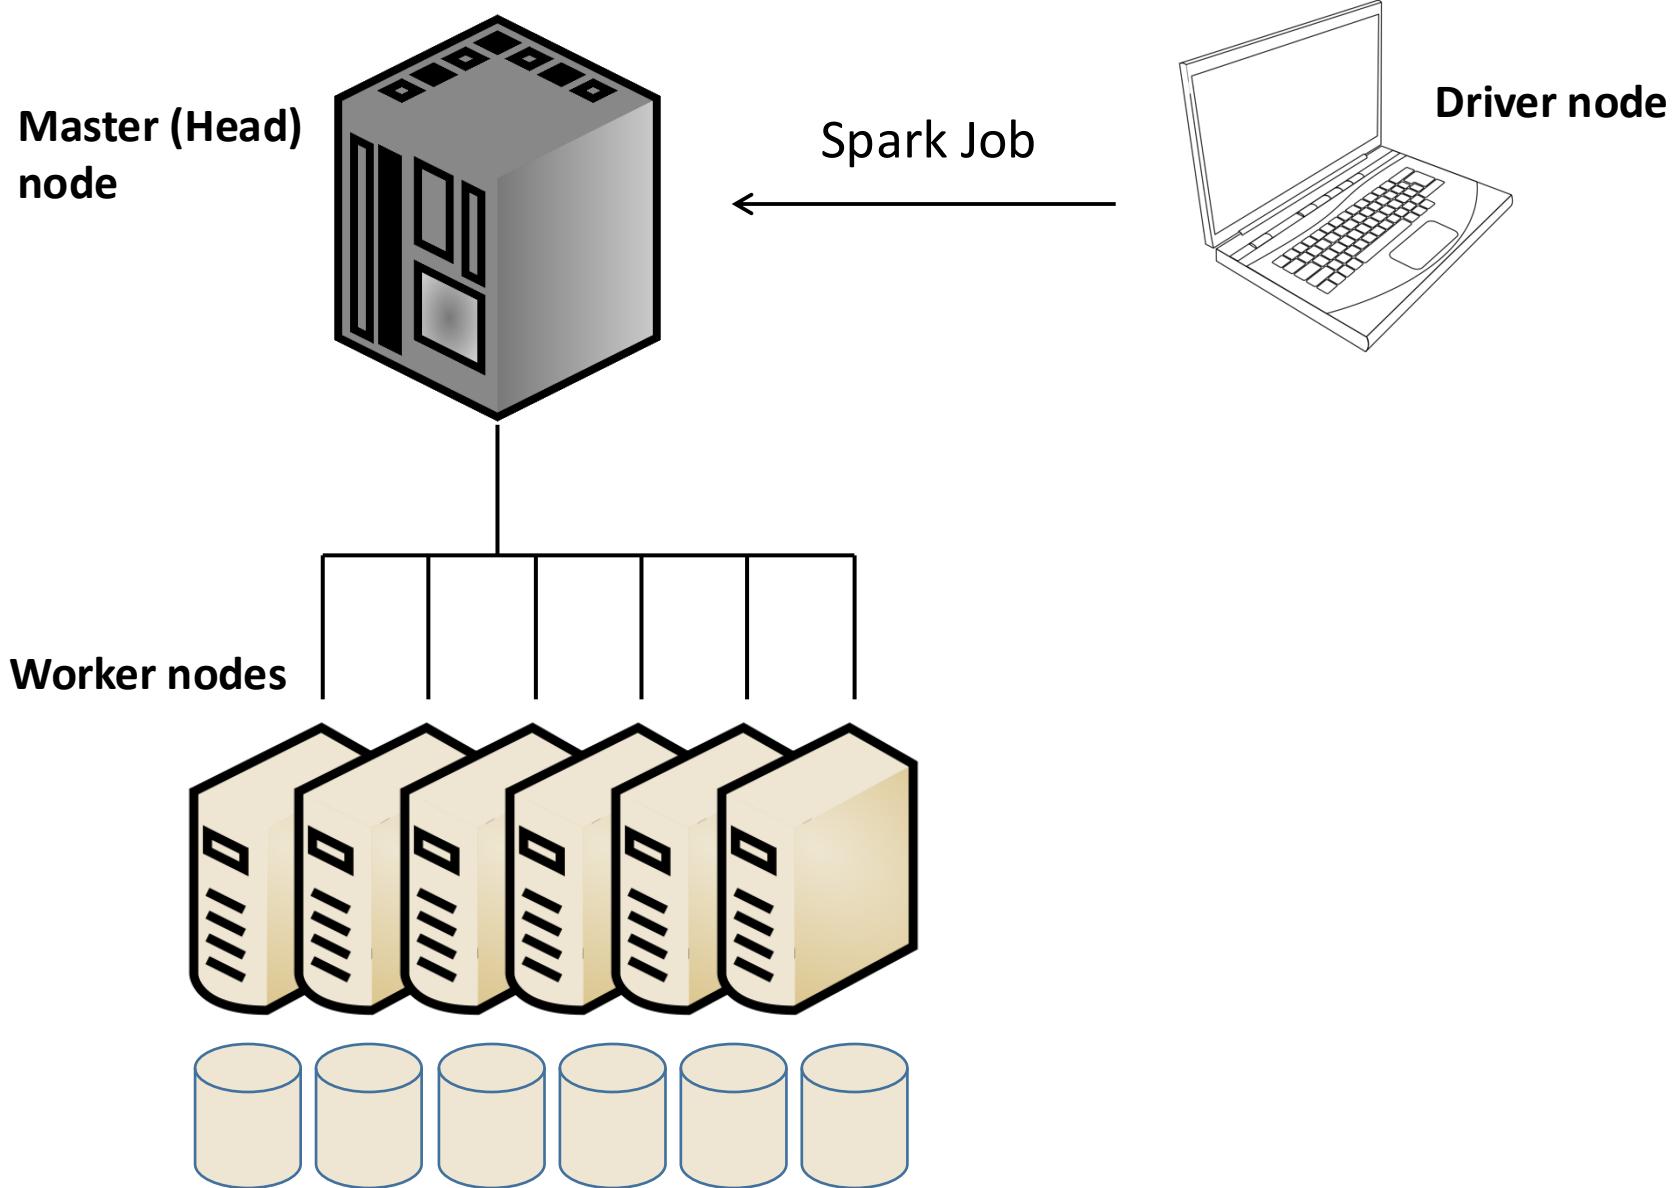
- A programming paradigm for expressing distributed algorithms
- Introduced by Google in 2004
  - Google File System for distributed storage
  - Google MapReduce for distributed processing
- Hadoop is the open-source counterpart released in 2007 and contributed mainly by Yahoo!
  - HDFS
  - Hadoop MapReduce



# Spark

- Hadoop and MapReduce were perfect as research vehicles
- They helped in framing what we really want in a big data system
- Spark came as a new system designed from scratch to satisfy the real need of big data
- A distributed shared-nothing system
- Uses a functional programming paradigm

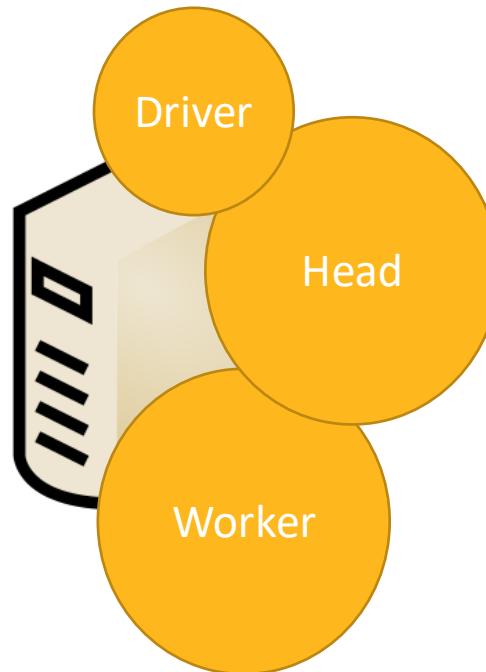
# Spark Overview



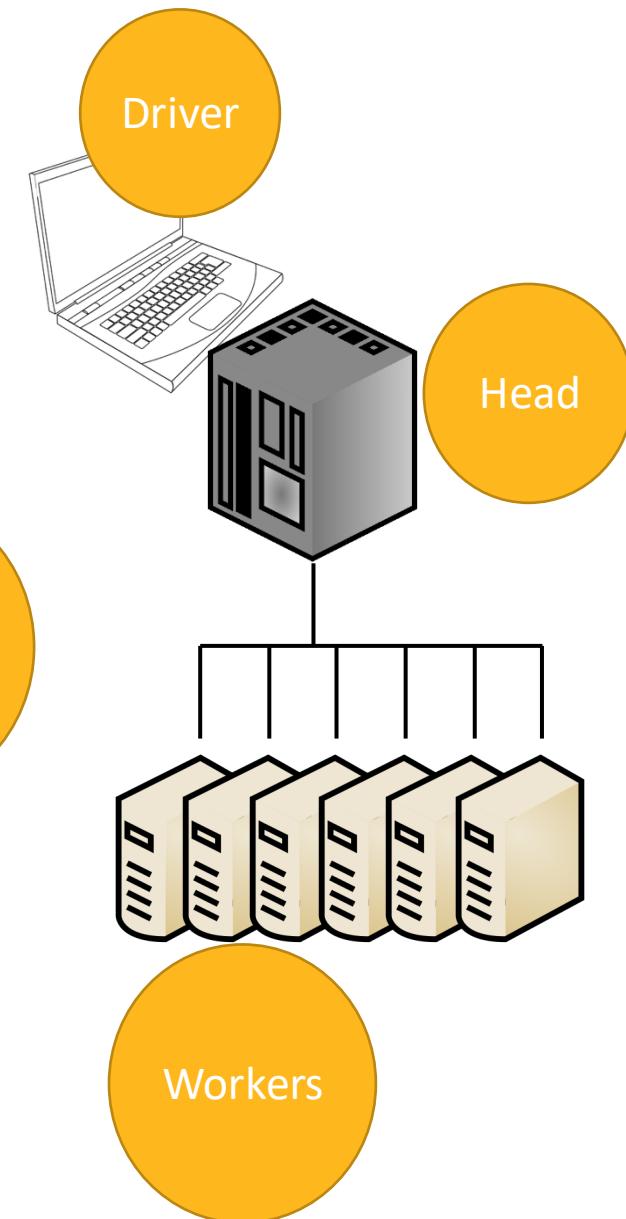
# Spark Operation Modes



Local mode



Stand-alone mode



Cluster mode

# Examples

# Examples

host	time	method	url	response	bytes
pppa006.compuserve.com	807256800	GET	/images/launch-logo.gif	200	1713
vcc7.langara.bc.ca	807256804	GET	/shuttle/missions/missions.html	200	8677

```
# Initialize the Spark context
JavaSparkContext spark =
    new JavaSparkContext("local", "CS226-Demo");
```

# Examples

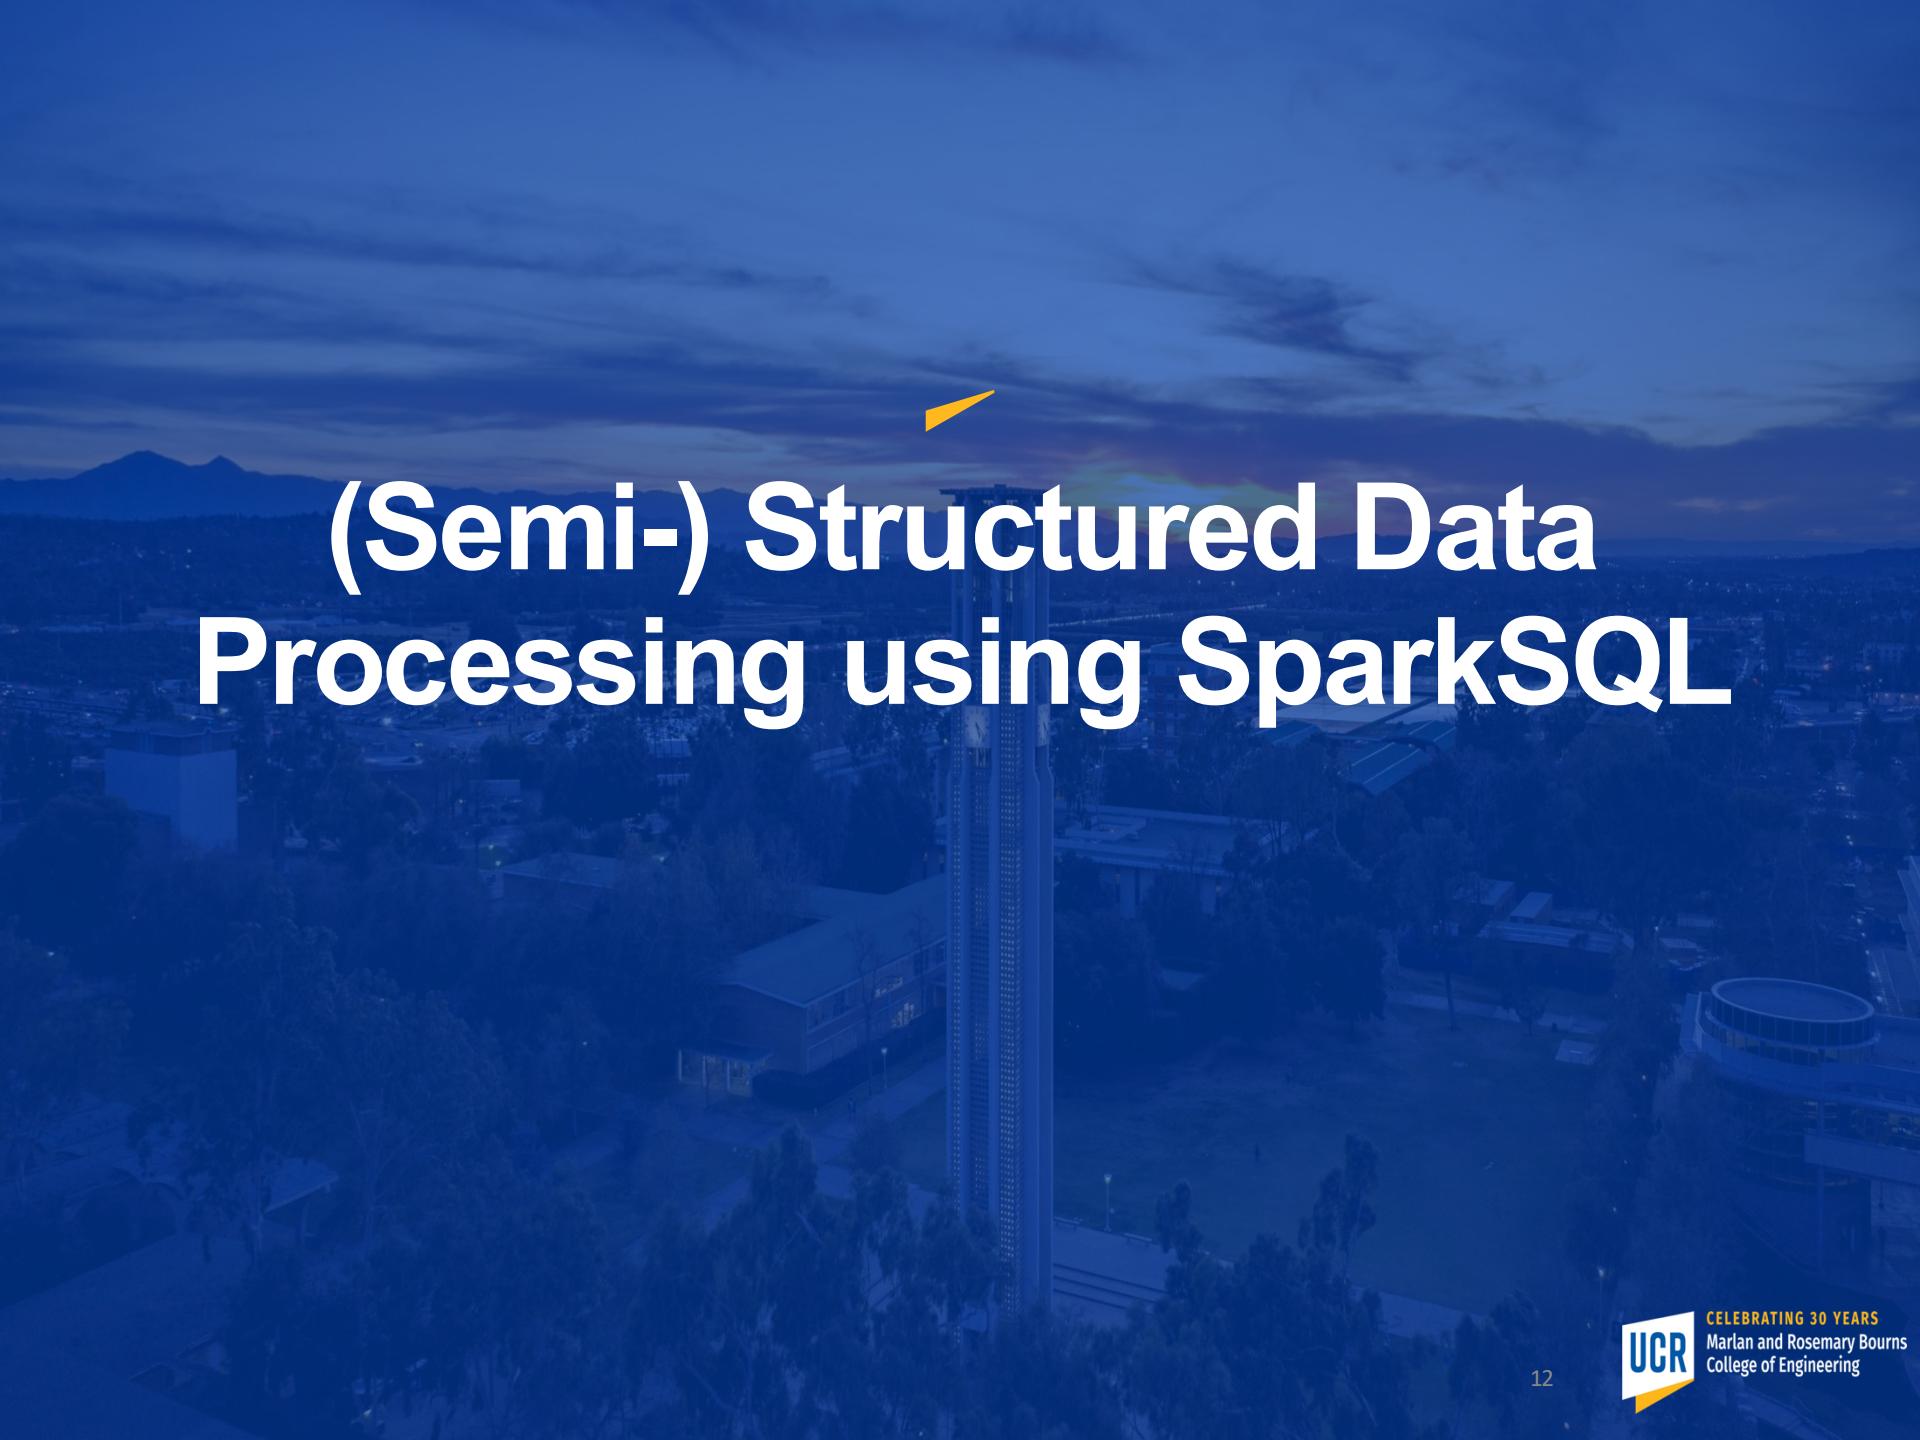
```
// Initialize the Spark context
JavaSparkContext spark =
    new JavaSparkContext("local", "CS226-Demo");

// Hello World! Example. Count the number of lines in the file
JavaRDD<String> textFileRDD =
    spark.textFile("nasa.tsv");
long count = textFileRDD.count();
System.out.println("Number of lines is "+count);
```

# Examples

```
// Count the number of OK lines
JavaRDD<String> okLines =
    textFileRDD.filter(s -> s.split("\t")[5].equals("200"));

long count = okLines.count();
System.out.println("Number of OK lines is "+count);
```



◆

# (Semi-) Structured Data Processing using SparkSQL

# Structured Data Processing

- A common use case in big-data is to process structured or semi-structured data
- In Spark RDD, all functions and objects are black-boxes.
- Any structure of the data has to be part of the functions which includes:
  - Parsing
  - Conversion
  - Processing

# SparkSQL

- Redesigned to consider Spark query model
- Supports all the popular relational operators
- Can be intermixed with RDD operations
- Uses the Dataframe API as an enhancement to the RDD API

Dataframe = RDD + schema

# Built-in operations in SprkSQL

- Filter (Selection)
- Select (Projection)
- Join
- GroupBy (Aggregation)
- Load/Store in various formats
- Cache
- Conversion between RDD (back and forth)

# SparkSQL Examples

# Code Setup

```
SparkSession sparkS = SparkSession  
    .builder()  
    .appName("Spark SQL examples")  
    .master("local")  
    .getOrCreate();
```

```
Dataset<Row> log_file = sparkS.read()  
    .option("delimiter", "\t")  
    .option("header", "true")  
    .option("inferSchema", "true")  
    .csv("nasa_log.tsv");  
  
log_file.show();
```

# Filter Example

```
// Select OK lines  
  
Dataset<Row> ok_lines =  
log_file.filter("response=200");  
long ok_count = ok_lines.count();  
System.out.println("Number of OK lines is  
"+ok_count);
```

```
// Grouped aggregation using SQL  
  
Dataset<Row> bytesPerCode =  
log_file.sqlContext().sql("SELECT response,  
sum(bytes) from log_lines GROUP BY response");
```



# MLlib: Machine learning in Spark

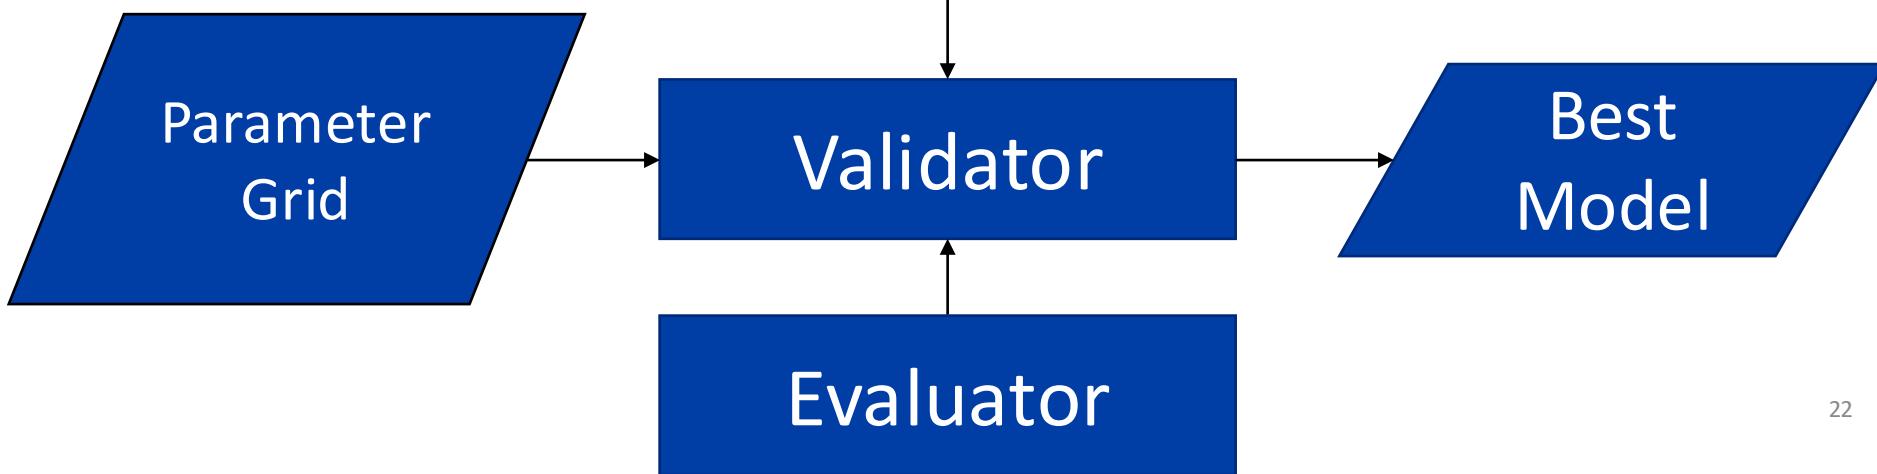
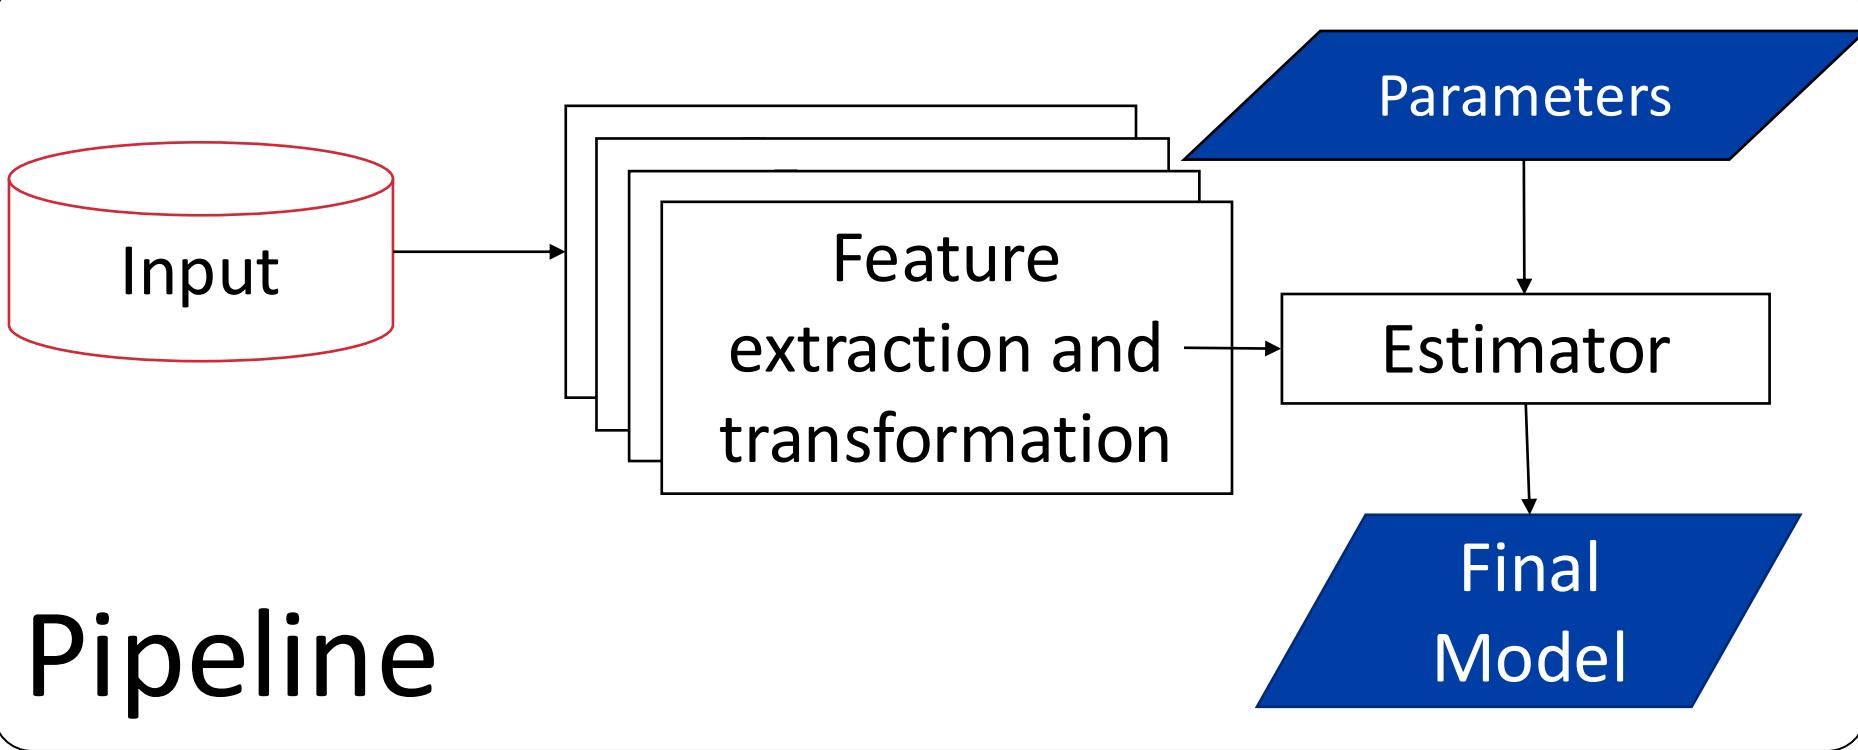
# Machine Learning Algorithms

- Supervised learning
  - Given a set of features and labels
  - Builds a model that predicts the label from the features
  - E.g., classification and regression
- Unsupervised learning
  - Given a set of features without labels
  - Finds interesting patterns or underlying structure
  - E.g., clustering and association mining

# Basic Statistics

- Column statistics
  - Minimum, Maximum, count, ... etc.
- Correlation
  - Pearson's and Spearman's correlation
- Hypothesis testing
  - Chi-square Test  $\chi^2$

# ML Pipeline



# Code Example on MLlib

# Input Data

House ID	Bedrooms	Area (sqft)	...	Price
1	2	1,200		\$200,000
2	3	3,200		\$350,000
...				

- Goal: Build a model that estimates the price given the house features, e.g., # of bedrooms and area

# Initialization

- Similar to SparkSQL

```
val spark = SparkSession  
  .builder()  
  .appName("SparkSQL Demo")  
  .config(conf)  
  .getOrCreate()
```

```
// Read the input  
val input = spark.read  
  .option("header", true)  
  .option("inferSchema", true)  
  .csv(inputfile)
```

# Transformations

```
// Create a feature vector  
val vectorAssembler = new VectorAssembler()  
.setInputCols(Array("bedrooms", "area"))  
.setOutputCol("features")
```

```
val linearRegression = new LinearRegression()  
.setFeaturesCol("features")  
.setLabelCol("price")  
.setMaxIter(1000)
```

# Create a Pipeline

```
val pipeline = new Pipeline()  
.setStages(Array(vectorAssembler, linearRegression))
```

// Hyper parameter tuning

```
val paramGrid = new ParamGridBuilder()  
.addGrid(linearRegression.regParam,  
         Array(0.3, 0.1, 0.01))  
.addGrid(linearRegression.elasticNetParam,  
         Array(0.0, 0.3, 0.8, 1.0))  
.build()
```

# Cross Validation

```
val crossValidator = new CrossValidator()  
  .setEstimator(pipeline)  
  .setEvaluator(new  
    RegressionEvaluator().setLabelCol("price"))  
  .setEstimatorParamMaps(paramGrid)  
  .setNumFolds(5)  
  .setParallelism(2)  
  
val Array(trainingData, testData) =  
  input.randomSplit(Array(0.8, 0.2))  
val model = crossValidator.fit(trainingData)
```

# Apply the model on test data

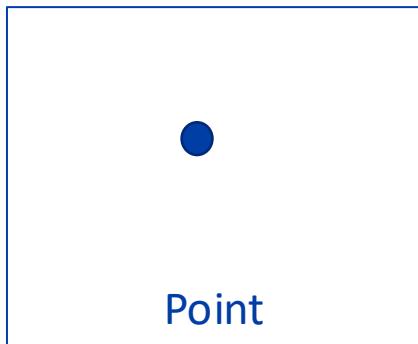
```
val predictions = model.transform(testData)
// Print the first few predictions
predictions.select("price", "prediction").show(5)
```

```
val rmse = new RegressionEvaluator()
.setLabelCol("price")
.setPredictionCol("prediction")
.setMetricName("rmse")
.evaluate(predictions)
println(s"RMSE on test set is $rmse")
```

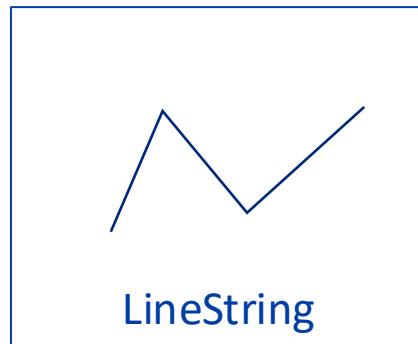


# Big Spatial Data Management on Spark

# Geometry Data Types



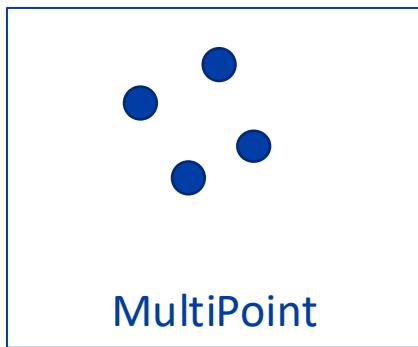
Point



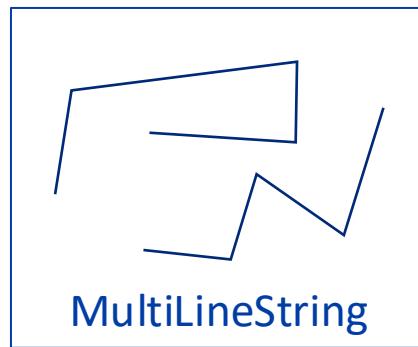
LineString



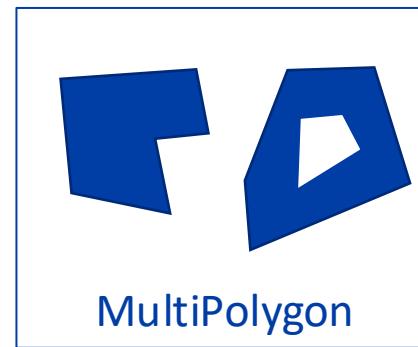
Polygon



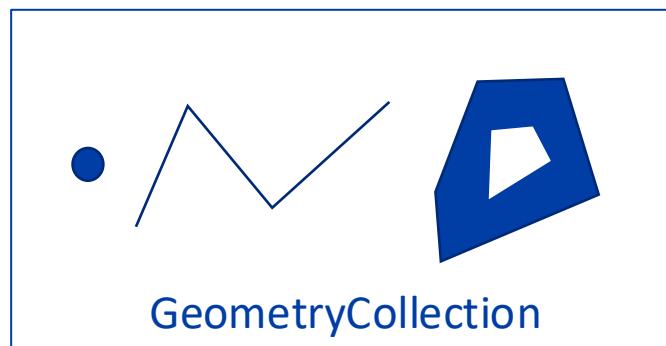
MultiPoint



MultiLineString

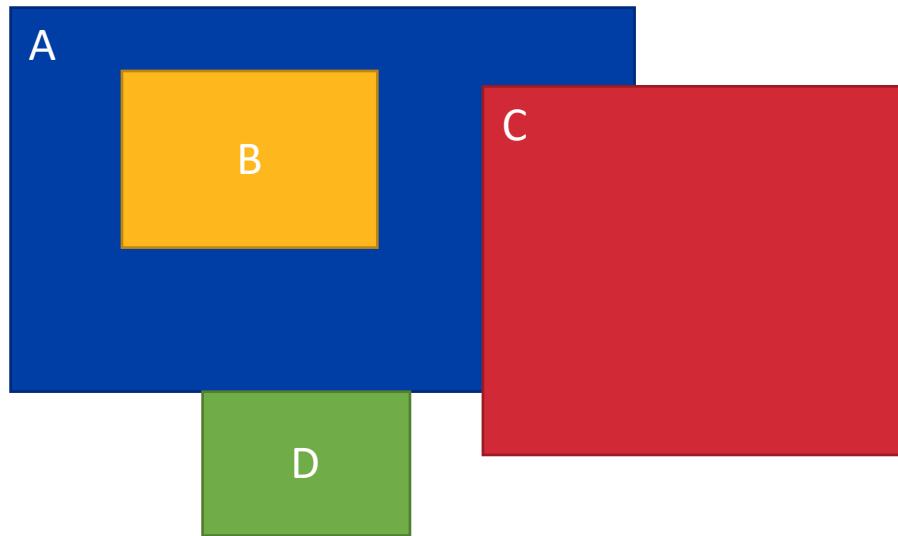


MultiPolygon



GeometryCollection

# Geometry Predicates



A Contains B  
A Overlaps C  
B Disjoint C  
A Touches D

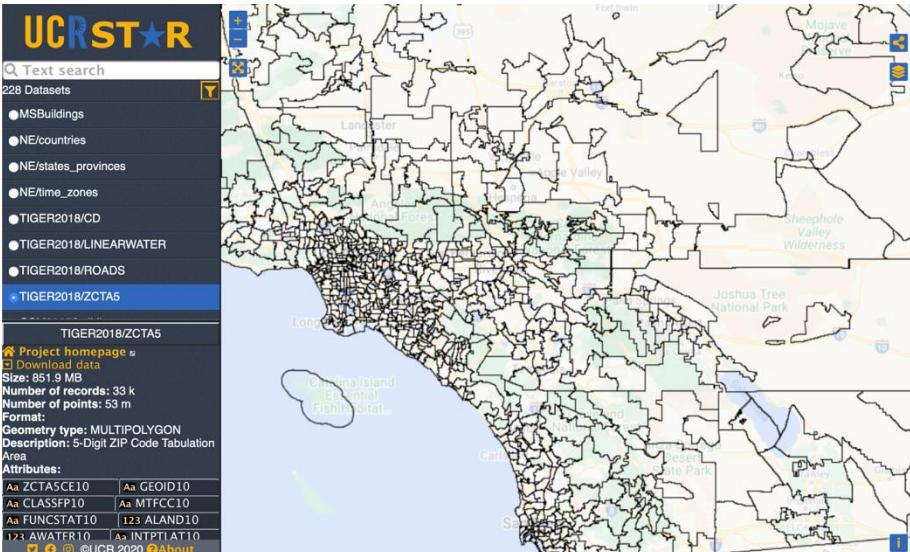
# Spatial Feature (IFeature)

Feature = Geometry + Other Attributes

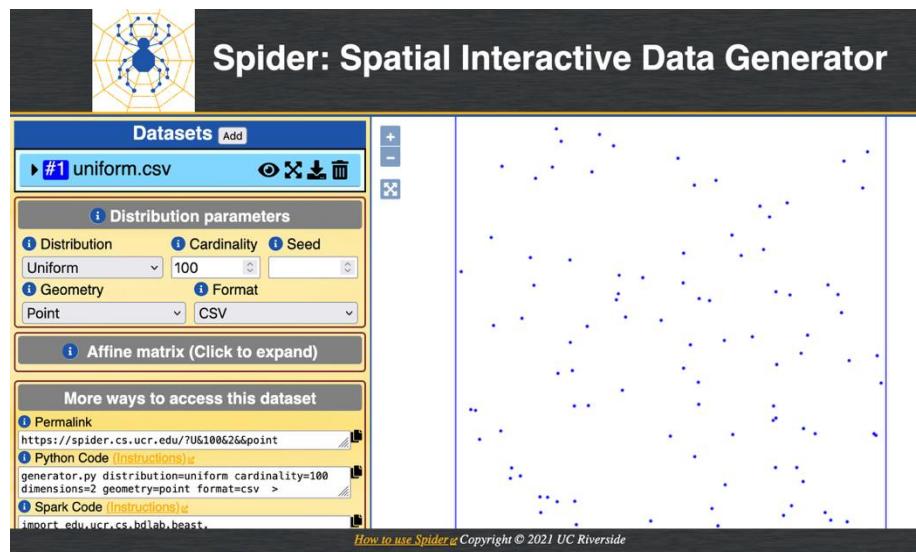
- Example
  - Road(Geometry, Name, Speed Limit)
  - State(Geometry, Name, Population)

# Data Source

- [UCRStar.com](http://UCRStar.com)
- 200+ datasets
- Full/subset download
- Standard formats



- [spider.cs.ucr.edu](http://spider.cs.ucr.edu)
- Data generator



# Data Loading

```
// Load a shapefile
val polygons: RDD[IFeature] =
sc.shapefile("tl_2018_us_state.zip")
// Load GeoJSON file
val points = sc.geojsonFile("Tweets.geojson")

// Load points from a CSV file
val lines = sc.readCSVPoint("Crimes.csv",
    "Longitude", "Latitude", ',', ',',
skipHeader = true)

// Load geometries from a CSV file
val lines = sc.readWKTFile("States.csv", 0,
'\t', skipHeader = false)
```

# Simple Manipulation

```
// Calculate the area and append as a new  
attribute  
polygons.map(f => Feature.append(f,  
    value = f.getGeometry.getArea, name = "area"))  
  
// Simplify the geometries into their convex hull  
polygons.map(f => {  
    val convex_hull = f.getGeometry.convexHull()  
    Feature.create(f, geometry = convex_hull)  
})
```

# Range Filters

```
// Select the geometry of the state of California
val california: IFeature = polygons.filter(f =>
f.getAttributeValue("NAME") == "California").first()

// Filter the points that are inside the state of
// California
val californiaPoints =
points.rangeQuery(california.getGeometry)
println(s"Number of points in California
${californiaPoints.count()}")
```

## Output

Number of points in California 259657

# Spatial Join

```
// Count points per state
val airportCountByState =
polygons.spatialJoin(airports)
  .map(fv => (fv._1.getAs[String]("NAME"), 1))
  .countByKey()

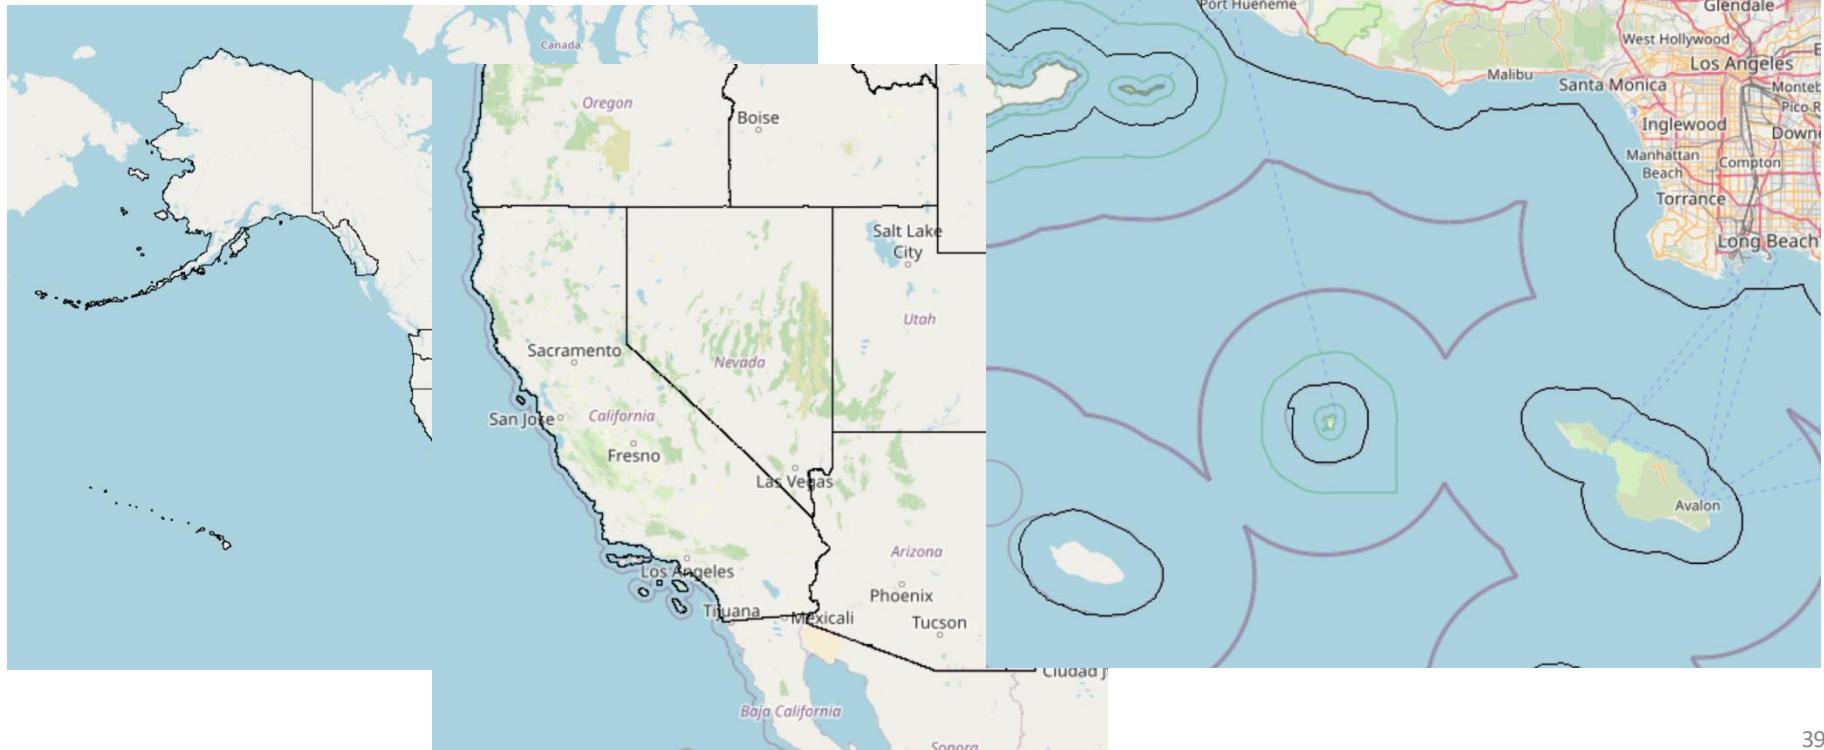
airportCountByState.foreach(sv =>
println(s"${sv._1}\t${sv._2}"))
```

## Output

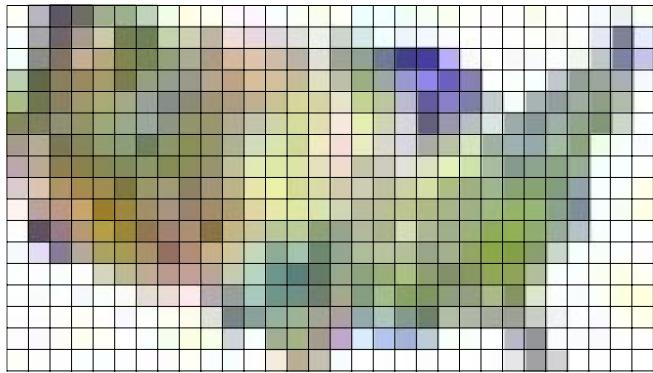
```
New Mexico      1
Connecticut1
Commonwealth of the Northern Mariana Islands    2
California     12
Nevada 3
```

# Visualization on a Map

```
// Plot states as a multilevel map  
polygons.plotPyramid("states", 10,  
    opts = "mercator" -> "true")
```



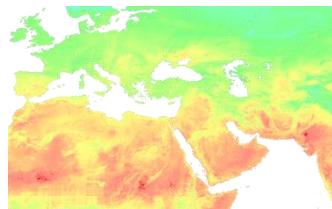
# Raster Data Representation



2D Array of values (pixels)



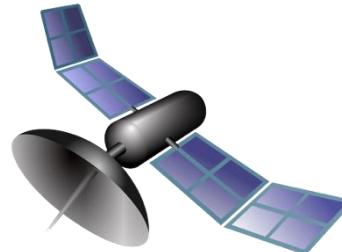
Vegetation



Temperature

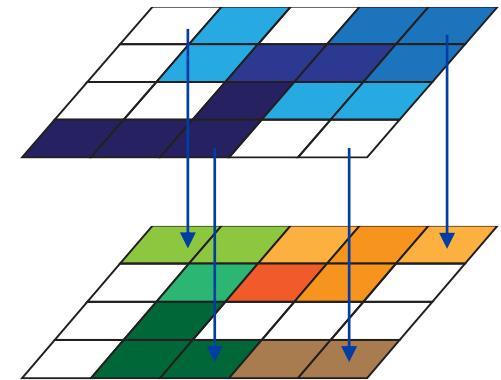


Camera



Satellite

Map Algebra



Linear algebra operations  
Map Algebra  
Overlay  
Rescale

# Raster (Satellite) Data Processing

```
// Load a raster file (or directory)
val raster: RDD[ITile[Int]] =
sc.geoTiff[Int]("glc2000_v1_1.tif")

// Load temperature in Kelvin from an HDF file
val temperatureK: RasterRDD[Float] =
sc.hdfFile("MOD11A1.A2022173.h08v05.006.2022174092443.hdf"
, "LST_Day_1km")

// Convert Kelvin to Fahrenheit
val temperatureF: RasterRDD[Float] =
temperatureK.mapPixels(k => (k-273.15f) * 9 / 5 + 32)

// Save the result as a GeoTIFF file
temperatureF.saveAsGeoTiff("temperature_f")
```

# Raster: Filter Pixels

```
val temperatureK: RasterRDD[Float] =  
  
sc.hdfFile("MOD11A1.A2022173.h08v05.006.2022174092443  
.hdf", "LST_Day_1km")  
  
// Keep only pixels with temperature > 300°K  
temperatureK.filterPixels(_>300)  
    .saveAsGeoTiff("temperature_high")
```

# Raster: Rescale and Reproject

```
val raster: RasterRDD[Int] =  
sc.geoTiff[Int]("glc2000_v1_1.tif")  
// Downscale a big raster to 360x180  
val rescaled = raster.rescale(360, 180)  
// Save as a single file  
rescaled.saveAsGeoTiff("glc_small",  
GeoTiffWriter.WriteMode -> "compatibility")
```

# Raster: Raster-Vector Join

```
// Load a raster file of Global Land Cover
val raster: RasterRDD[Int] =
sc.geoTiff[Int]("glc2000_v1_1.tif")

// Filter the pixels that represent trees
val trees = raster.filterPixels(lc => lc >= 1 && lc <= 10)

// Load all countries
val countries =
sc.shapefile("ne_10m_admin_0_countries.zip")

// Count number of tree pixels per country
val result = trees.raptorJoin(countries)
.map(x => x.feature.getAs[String]("NAME"))
.countByValue().toMap
```



**Big Data Management & Analysis**

**with**

**Apache AsterixDB**

# AsterixDB Overview

# AsterixDB: “One Size Fits a Bunch!”

## *Wish-list:*

- Able to **manage** data
- **Flexible** data model
- Full **query** capability
- Continuous data **ingestion**
- Efficient and robust **parallel** runtime
- Cost **proportional** to task at hand
- Support today’s “**Big Data**” data types”

Semistructured  
data management

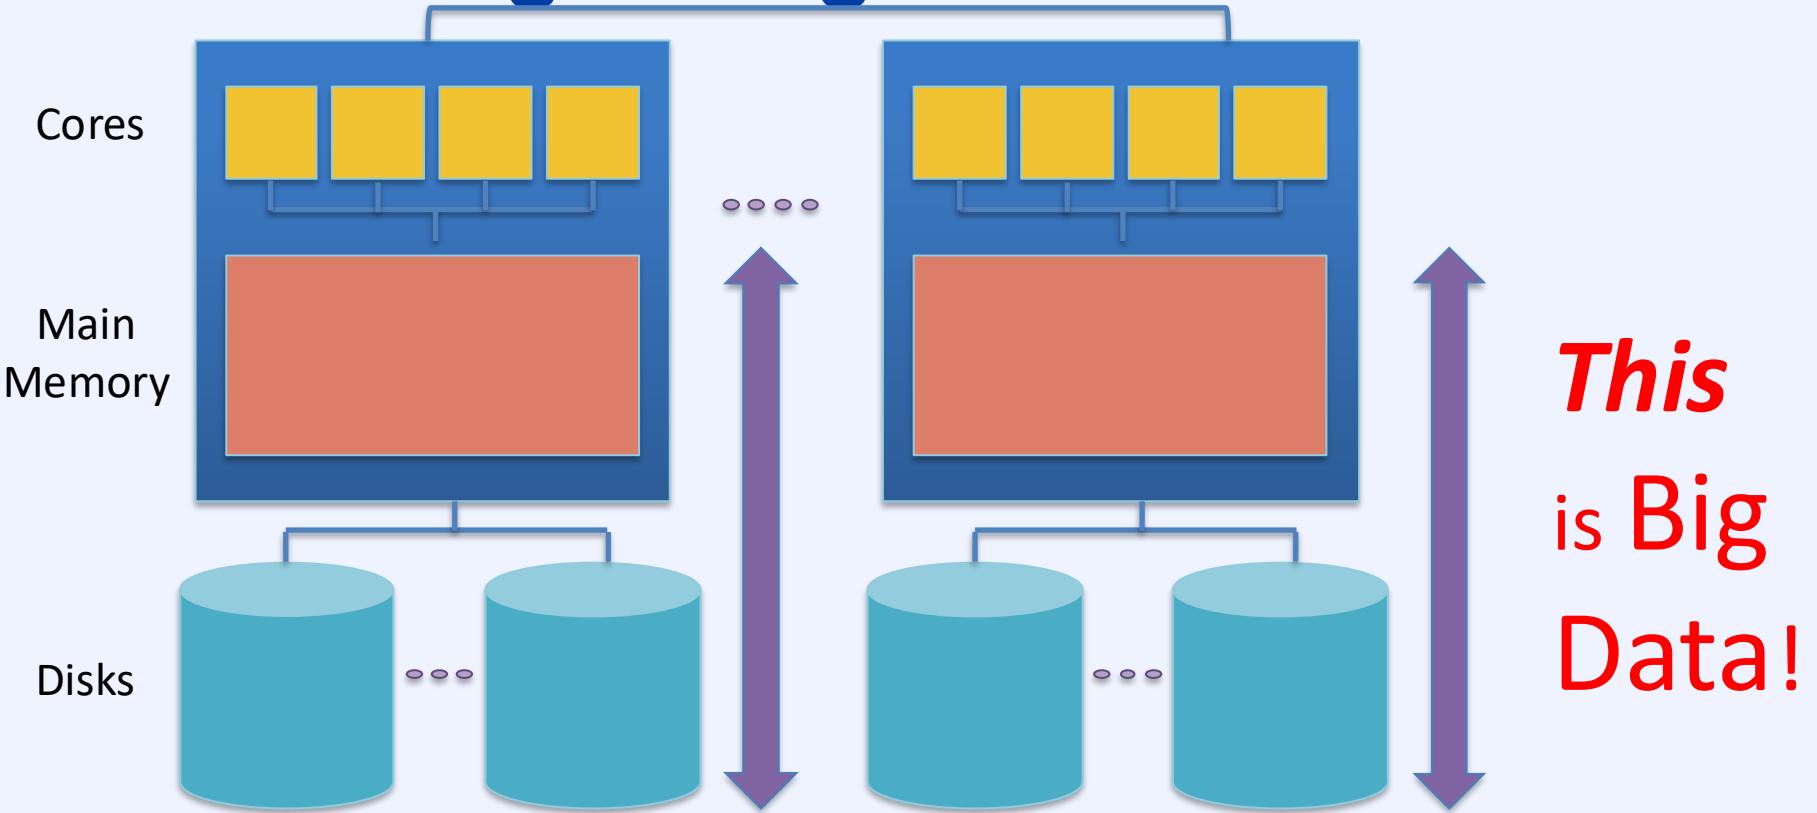


Parallel  
DB systems

First-gen BD  
analysis tools

→ *Parallel NoSQL DBMS* ←

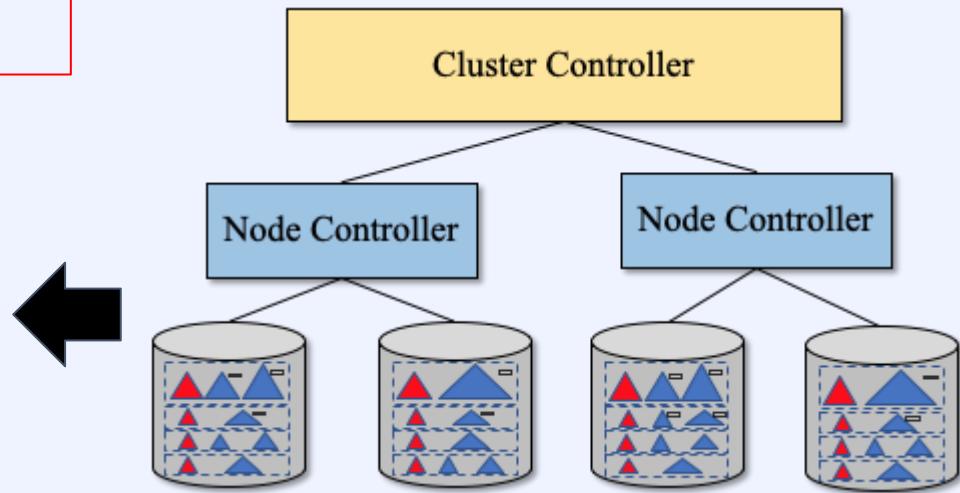
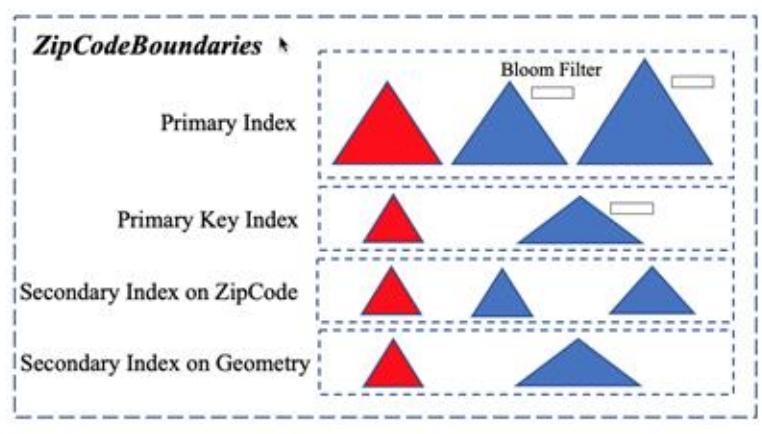
# Just How Big is “Big Data”?



# An Indexed Analytics Dataset

## Partitioned local storage approach

- Hashed on primary key (PK)
- Primary index w/ PK + record
- Secondary index(es) with SK + PK
- Record updates are always local



# Spatial Data Types and Functions

# AsterixDB User Interface (<http://localhost:19006>)

The screenshot shows the AsterixDB User Interface. On the left, the 'QUERY INPUT' panel displays a single query entry '1'. Below it, 'WARNINGS(0)' and a 'CLEAR' button are visible. In the center, there's an 'EXPLAIN' button followed by navigation arrows. On the right, the 'METADATA INSPECTOR' panel is open, showing a sidebar with 'DATAVERSES' (Default, GeospatialData, Metadata, SampleGeospatialData), 'REFRESH' button, and sections for 'DATASETS', 'DATATYPES', 'INDEX', and 'USER DEFINED FUNCTIONS'.

WEBSITE FILE ISSUES DOCUMENTATION CONTACT GITHUB

QUERY INPUT (1/1)

Default PLAN FORMAT OUTPUT FORMAT

Default JSON JSON

QUERY HISTORY

1

WARNINGS(0)

CLEAR EXPLAIN ▶

METADATA INSPECTOR

DATAVERSES

- Default
- GeospatialData
- Metadata
- SampleGeospatialData

REFRESH

DATASETS

DATATYPES

INDEX

USER DEFINED FUNCTIONS

51

# Creating a Dataverse - Query 1

```
DROP DATAVERSE UNIQUE_DATaverse_NAME IF EXISTS;  
CREATE DATAVERSE UNIQUE_DATaverse_NAME IF NOT EXISTS;
```

Please do not forget to replace *UNIQUE\_DATaverse\_NAME* with your

The screenshot shows the AsterixDB query interface. On the left, the 'QUERY INPUT (1/1)' panel displays a single line of code: '1 USE GeospatialDatasets;'. A red arrow points from the text 'Please do not forget to replace UNIQUE\_DATaverse\_NAME with your' to the 'GeospatialDatasets' dropdown in this panel. The 'OUTPUT FORMAT' dropdown is set to 'JSON'. On the right, the 'METADATA INSPECTOR' panel lists 'DATAVERSES' with three options: 'Default', 'GeospatialDatasets', and 'Metadata'. The 'Default' option is selected. Below it are sections for 'DATASETS', 'DATATYPES', 'INDEX', and 'USER DEFINED FUNCTIONS'. At the bottom of the interface, there are buttons for 'EXPLAIN', 'REFRESH', and navigation arrows.

You need to select your dataset before running any query.

# Loading Data to AsterixDB

```
// Chicago Crimes Dataset
CREATE TYPE ChicagoCrimesType IF NOT EXISTS AS {
    id: uuid,
    g: geometry,
    `Primary Type`: String
};

CREATE DATASET ChicagoCrimes(ChicagoCrimesType) IF
NOT EXISTS PRIMARY KEY id AUTOGENERATED;

LOAD DATASET ChicagoCrimes USING localfs
(("path"="1:///home/admin/bosdata/chicagocrimes.js
on"), ("format"="adm"));

SELECT VALUE n
FROM ChicagoCrimes n LIMIT 1;
```

```
{
    "ChicagoCrimes": {
        "id": "d85704ce-ac9f-65f0-d6d6-2deb29206d07",
        "g": {
            "type": "Point",
            "coordinates": [
                -87.883611316,
                41.980826277
            ],
            "crs": {
                "type": "name",
                "properties": {
                    "name": "EPSG:4326"
                }
            }
        },
        "ID": "9805746",
        "Case Number": "HX442584",
        "Primary Type": "MOTOR VEHICLE THEFT",
        "Description": "AUTOMOBILE",
        "Location Description": "AIRPORT VENDING
ESTABLISHMENT",
        "Arrest": "false",
        "Domestic": "false",
        ...
        "Updated On": "02/10/2018 03:50:01 PM"
    }
}
```

# External Dataset

```
// External datasets
CREATE TYPE RoadsTypeExternal IF NOT EXISTS AS {
    g: geometry,
    FULLNAME: String
};

CREATE EXTERNAL DATASET RoadsExternal(RoadsTypeExternal)
USING localfs
(("path"="127.0.0.1:///home/admin/bossdata/roads.json"),
("format"="adm"));

SELECT VALUE r
FROM RoadsExternal r
LIMIT 1;
```

Good for in-situ processing of static data (Similar to Hadoop and Spark)

# Example Query 12

**Problem:** List the Roads that cross each other.

We do a self  
join

```
SELECT r1.FULLNAME AS r1_name, r2.FULLNAME AS r2_name  
FROM Roads r1, Roads r2  
WHERE st_crosses(r1.g, r2.g) AND r2.id > r1.id  
ORDER BY r1.FULLNAME;
```

We check if two roads are crossing each  
other and apply a duplicate avoidance.

# Multi-dataset and nested queries - Query 13

**Problem:** Find the average number of neighbors for Zip Code boundaries.

- Do a self join on `ZipCodeBoundaries`
- Check if Geometries touch or not with `st_touches()`
- Use GROUP BY `ZCTA5CE10` - the zip code value
- Use COUNT()
- Use AVG()

```
SELECT AVG(c) FROM (
    SELECT COUNT(1) as c
        FROM ZipCodeBoundaries AS z1, ZipCodeBoundaries AS z2
       WHERE st_touches(z1.g, z2.g)
      GROUP BY z1.ZCTA5CE10
) as t;
```



# Graph Processing

# Big Graphs

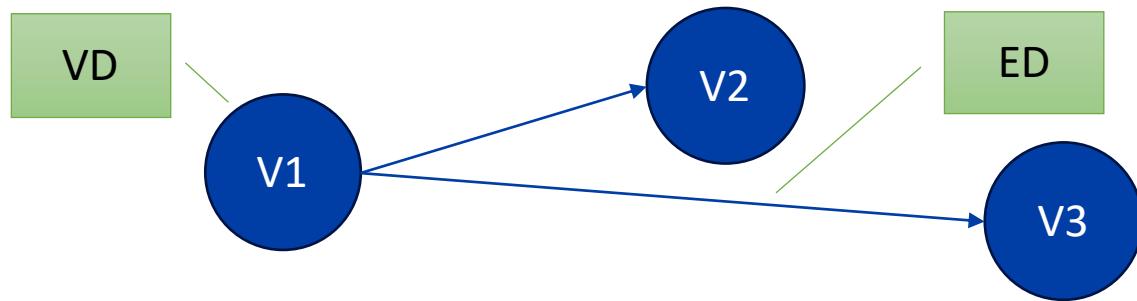
- Graphs with hundreds of millions of vertices and billions of edges

1.7 billion users  
100s of billions of connections

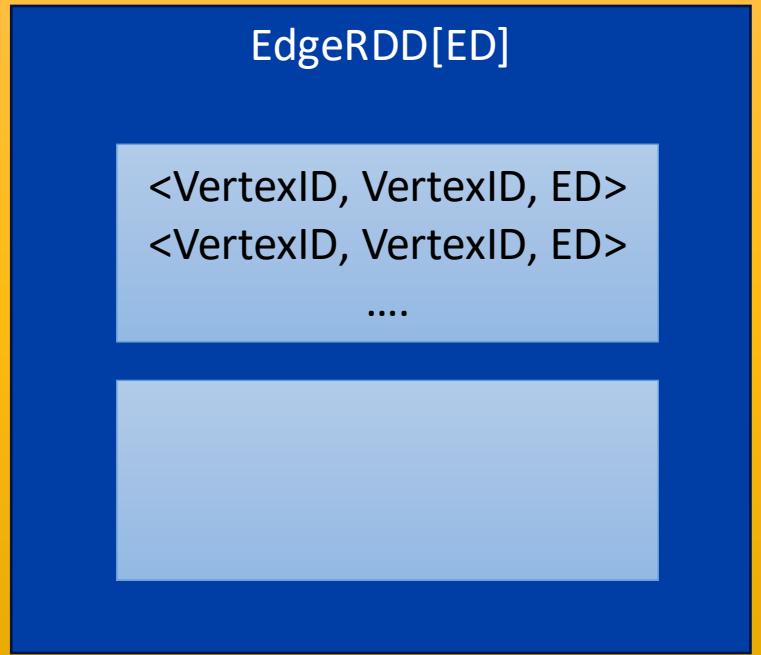


- The graph is too big to fit on one machine

# Property Graph Model



Graph[VD, ED]

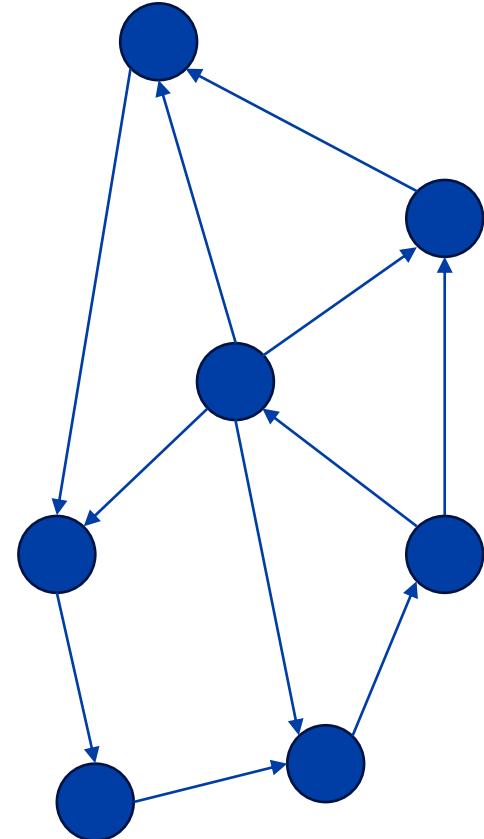


# Examples of Graphs

- Knowledge Graph
  - Graph[String, String]
  - VD: String – entity names
  - ED: String – relationship
- Road Network
  - VD: (longitude, Latitude)
  - ED: (StreetName, SpeedLimit, ...)

# Iterative Graph Processing

- Each vertex send a message along outgoing edges
- Each vertex updates its vertex data by aggregating incoming messages
- Repeat until a stopping condition
- Pregel API





# Vector DBMS

# The Rise of AI

- Unstructured queries are becoming more prevalent with the rise of Large Language Models (LLMs)
  - “Find the five cheapest 24” computer monitors that have at least 4.5 star

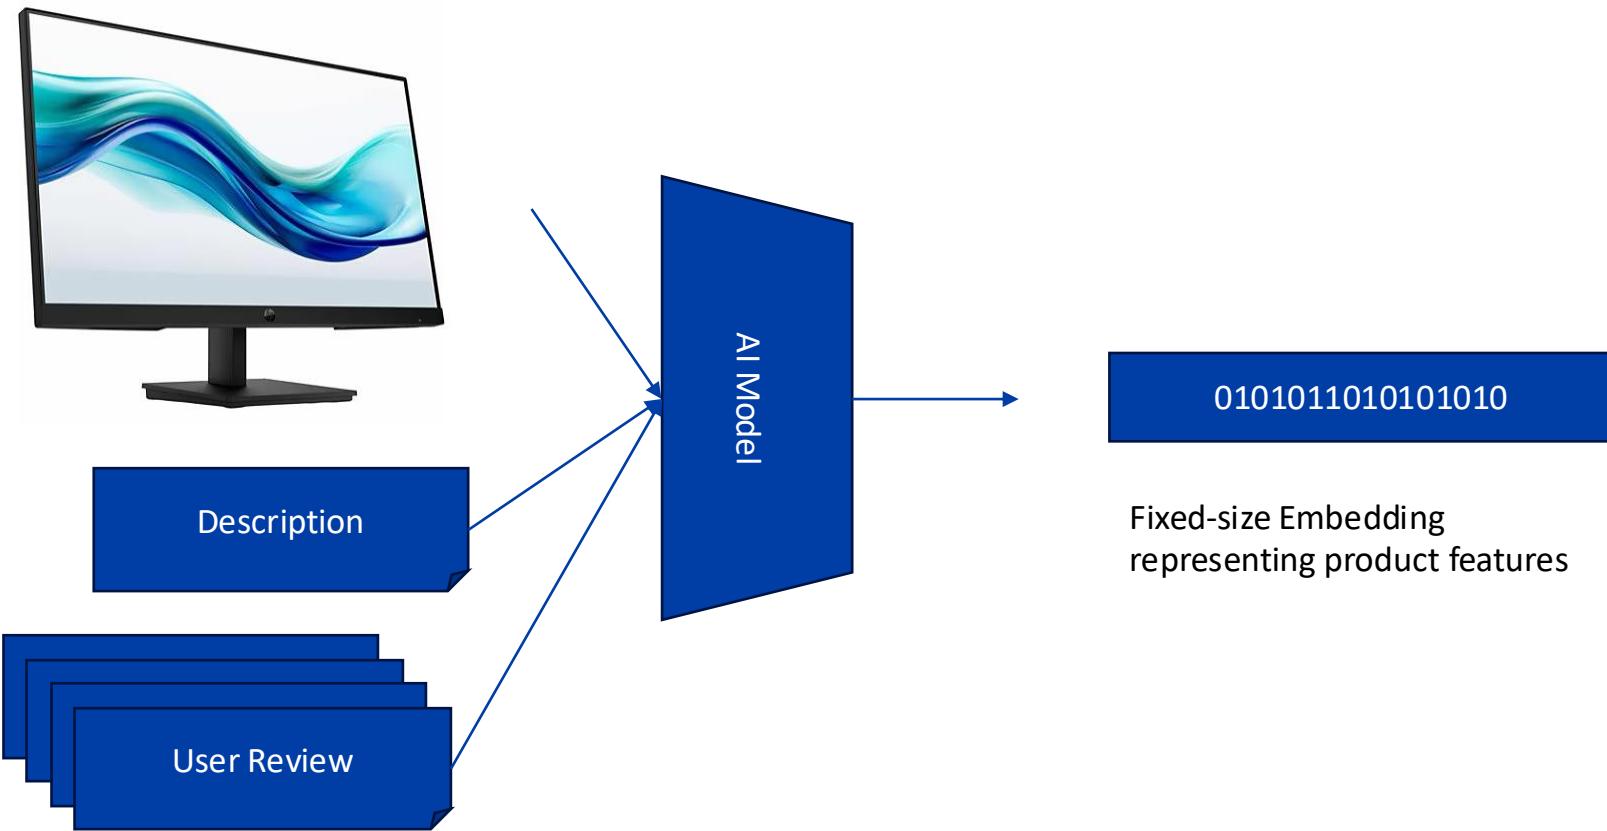
```
SELECT id, name, size_inches, price, rating
FROM monitors
WHERE size_inches = 24 AND rating >= 4.5
ORDER BY price ASC
LIMIT 5;
```

# More Interesting Unstructured Queries

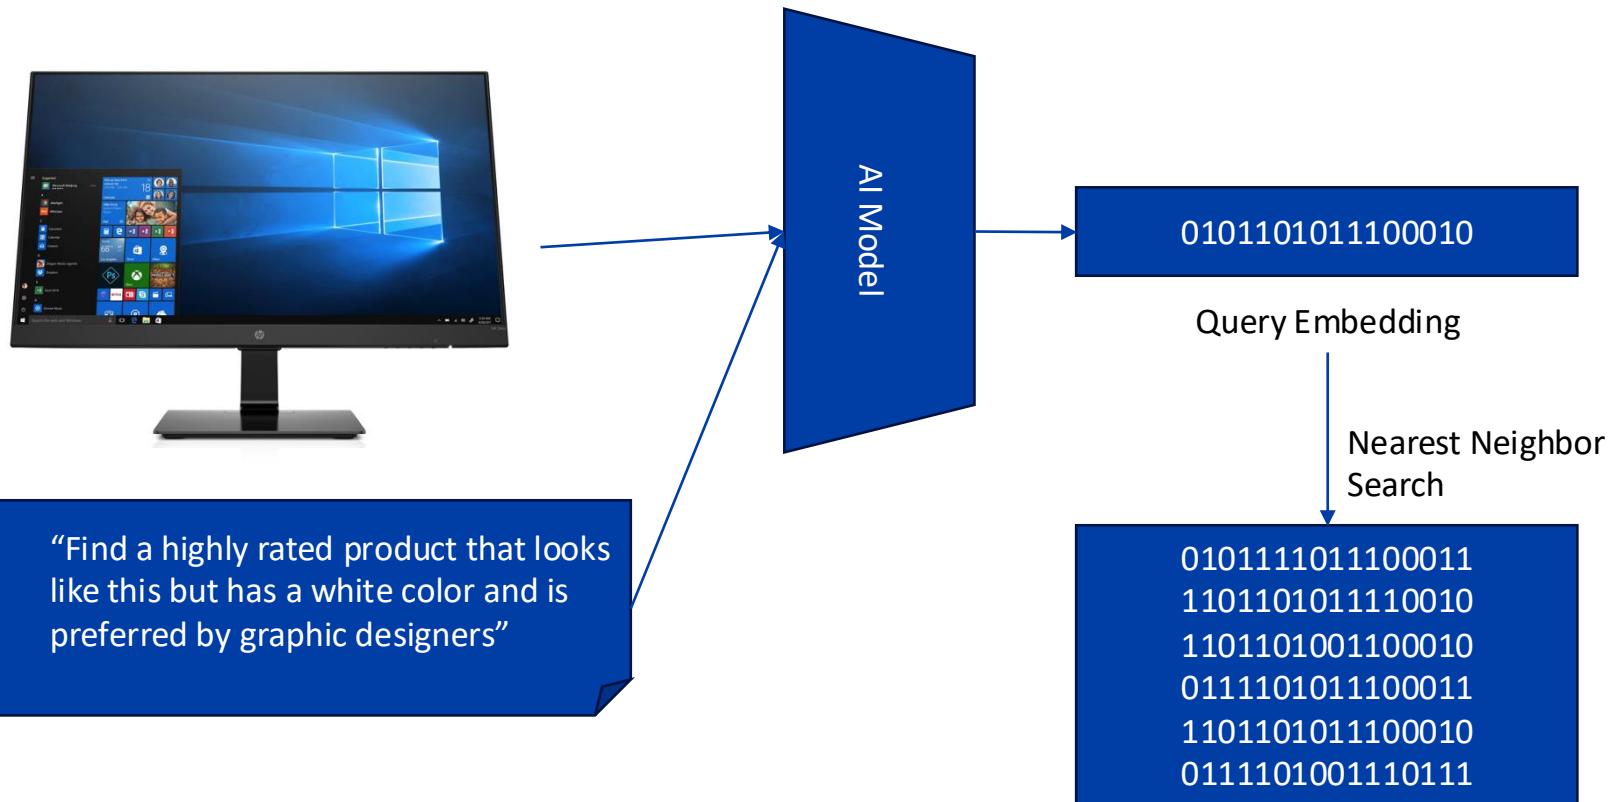
- In a database of products with pictures, description, and user reviews
  - “Find highly rated products that look like this but has a white color and is preferred by graphic designers”



# AI Embedding of Records



# Semantic Search



# RAG Question Answering

- RAG: Retrieval Augmented Generation
- Given a query:
  - Create a query embedding
  - Run an approximate nearest neighbor (ANN) query to find the most relevant documents
  - With an appropriate prompt, feed the documents through a generative AI model to get the final answer

