

# Big-data Storage

HDFS

# HDFS Overview




- A distributed file system
- Built on the architecture of Google File System (GFS)
- Shares a similar architecture to many other common distributed storage engines such as Amazon S3 and Microsoft Azure
- HDFS is a stand-alone storage engine and can be used in isolation of the query processing engine
- Even if you do not use Hadoop MapReduce, you will probably still use HDFS

# HDFS Topics

- HDFS design and architecture
- Fault tolerance in HDFS
- Command-line interface (HDFS Shell)
- Java API
- Create (Write) a file
- Stream reading
- Structured reading

# Background on Disk Storage

- What are file systems and why do we need them?
- A file is a logical sequence of bits/bytes
- A physical disk stores data in sectors, tracks, tapes, blocks, ... etc.




# File System

- Any file system, is a method to provide a high-level abstraction on physical disk to make it easier to store files






Files



Folders

File System



# Distributed File System









Files




Folders

Distributed File System




# HDFS Architecture






# What is where?

File and directory names  
Block ordering and locations  
Capacity of data nodes  
Architecture of data nodes




**Name node**




# Data Loading

The most common  
replication factor is three




# HDFS Storage




# Analogy to Unix FS

The logical view is similar




# Analogy to Unix FS

The physical model is comparable



Unix



HDFS



# Fault Tolerance in HDFS

# Replication


- The default fault tolerance mechanism in HDFS is replication
- The most common replication factor is three
- If one or two nodes are temporarily unavailable, the data is still accessible
- If one or two nodes permanently fail, the master node replicates the under-replicated blocks to reach the desired replication factor
- Drawback: reduced disk capacity

# Replication Overhead

- Example: A cluster with 12 nodes, each node has a 10TB disk, what is the actual HDFS capacity?
- Total physical disk available  
 $= 12 \times 10 = 120TB$
- Actual perceived capacity of HDFS =  
 $\frac{120}{3} = 40 \text{ TB}$
- The overhead is 200%

# Advanced Feature: Erasure Coding

- Uses advanced algorithms for recovery, e.g., Reed-Solomon, XOR



# How Redundancy Works

- Consider that we want to store the two values  $x$  and  $y$
- Instead of repeating the raw values, we can store the following four values:
  - $x, y, x + y, x - y$
- If we lose one or two values, we can still recover the other two by solving an array of linear equations (two in this case)
- This idea can be generalized with more efficient formulas that lowers the computation

# Storage Overhead with Erasure Coding

- HDFS Erasure Coding: (6, 3) RS Code
- $k = 6$  data blocks
- $m = 3$  parity blocks
- Storage overhead =  $\frac{6+3}{6} - 1 = 0.5$



# **HDFS Shell**

## **Command-line Interface (CLI)**

# HDFS Shell

- Used for common operations
- Its usage is similar to Unix shell commands
- Basic operations include
  - ls, cp, mv, mkdir, rm, ...
- HDFS-specific operations include
  - copyToLocal, copyFromLocal, setrep, appendToFile, ...

# HDFS Shell

- General format

```
hdfs dfs -<cmd> <arguments>
```

- So, instead of


```
mkdir -p myproject/mydir
```

- You will write

```
hdfs dfs -mkdir -p myproject/mydir
```

- A list of shell commands with usage
  - <https://hadoop.apache.org/docs/r3.3.0/hadoop-project-dist/hadoop-common/FileSystemShell.html>

# HDFS API



# HDFS API Classes

- Configuration: Holds system configuration such as where the master node is running and default system parameters, e.g., replication factor and block size
- Path: Stores a path to a file or directory
- FileSystem: An abstract class for file system operations

# Fully Qualified Path

**hdfs://namenode:9000/path/to/file**

hdfs: the file system scheme. Other possible values are file, ftp, s3, ...

namenode: the name or IP address of the node that hosts the the file system

9000: the port on which the master node is listening

/path/to/file: the absolute path of the file

# Shorter Path Forms

- “file”: relative path to the current working directory in the default file system
- “/path/to/file”: Absolute path to a file in the default\* file system (as configured)
- “hdfs:///path/to/file”: Use the default\* values for the master node and port
- “hdfs://masternode/path/tofile”: Use the given masternode name or IP and the default\* port
- “file:///home/user/file”: an absolute path in the local file system

\*All the defaults are in the Configuration object

# HDFS API (Java)

Create the file system

```
Configuration conf = new Configuration();
Path path = new Path("...");
FileSystem fs = path.getFileSystem(conf);

// To get the local FS
fs = FileSystem.getLocal (conf);

// To get the default FS
fs = FileSystem.get(conf);
```

# HDFS API

Create a new file

```
FSDataOutputStream out =  
fs.create(path, ...);
```

Delete a file

```
fs.delete(path, recursive);  
fs.deleteOnExit(path);
```

Rename/Move a file

```
fs.rename(oldPath, newPath);
```

# HDFS API

Open a file

```
FSDataInputStream in = fs.open(path, ...);
```

Seek to a different location

```
in.seek(pos);  
in.seekToNewSource(pos);
```

# HDFS API

Concatenate

```
fs.concat(destination, src[]);
```

Get file metadata

```
fs.getFileStatus(path);
```

Get block locations

```
fs.getFileBlockLocations(path, from, to);
```




# **Internals of HDFS**

## **Writing**

# Data Loading



The most common  
replication factor is three




# Physical Cluster Layout




# Analogy of racks




# HDFS Writing Process




# HDFS Writing Process




# HDFS Writing Process




# HDFS Writing Process




# HDFS Writing Process




# HDFS Writing Process



# HDFS Writing Process



# HDFS Writing Process




# Notes about writing to HDFS

- Data transfers of replicas are pipelined
- The data does **not** go through the name node
- Random writing is **not** supported
- Appending to a file is supported but it creates a new block

# Writing from a datanode

If the file creator is running on one of the data nodes, the first replica is always assigned to that node  
The second and third replicas are assigned as before






# **Internals of HDFS**

## **Stream reading**


# Reading from HDFS

- Reading is relatively easier
- No replication is needed
- Replication can be exploited
- Random reading *is* allowed


# HDFS Reading Process




# HDFS Reading Process




# HDFS Reading Process



# HDFS Reading Process




# HDFS Reading Process



# Reading from a datanode

1. If the block is locally stored on the reader, this replica is chosen to read
2. If not, a replica on another machine in the same rack is chosen
3. Any other random block replica is chosen



When self-reading occurs, HDFS can make it much faster through a feature called short-circuit

# Notes About Reading

- The API is much richer than the simple open/seek/close API
  - You can retrieve block locations
  - You can choose a specific replica to read
- The same API is generalized to other file systems including the local FS and S3
- Review question: Compare random access read in local file systems to HDFS




# Special Features in HDFS


# HDFS Special Features

- Node decommission
- Load balancer
- Cheap concatenation


# Node Decommission



# Load Balancing




# Load Balancing




Start the load balancer

# Cheap Concatenation



Name node



Concatenate File 1 + File 2 + File 3 → File 4

Rather than creating new blocks, HDFS can just change the metadata in the name node to delete File 1, File 2, and File 3, and assign their blocks to a new File 4 in the right order.



# **Internals of HDFS**

## **Structured Reading**

# If you start something, finish it



# Data files

- In distributed big data processing, input files contain records not just raw bytes
- We need to a way to read records from files in HDFS
- For efficiency, we should split the file and read it in parallel

# Block

- Part of the file that is **physically** stored in one or more data nodes

# Split

- A **logical** part of the file defined by an offset and length
- Might or might not align with blocks

# File Splitting

- How to split the file?
  - By record: For fixed-size records
  - By size: For variable-size records
- Considerations, splitting the file should be fast
- Should not need to read the entire file to split it

Input File

# Simple File Splitting

Block 1

Block 2

Block 3

Block 4

Input File

- A split is created for each block
- Advantages
  - Data locality
  - Efficiency
- Drawback
  - A record might span more than one split

# Simple File Splitting

Block 1

Block 2

Block 3

Block 4

Input File

- A split is created for each block
- Advantages
  - Data locality
  - Efficiency
- Drawback
  - A record might span more than one split

# Read data in every split

Split 1



- Which records will be read for each of the four splits?


Split 3

Split 4

# Reading process

- Split the file based on the file metadata
  - File size, block sizes, # of nodes
- Each split is defined by:
  - File name, Start offset, Length
- For each split:
  - Seek to the start offset
  - Skip the first record (except for the first split)
  - Read until the beginning of the record goes beyond the start + length

# Reading block-compressed files



# Conclusion

- HDFS is a general-purpose distributed file system
- It provides a write-once ready-many access interface
- Supports random reading which can be used for stream reading and structured reading
- Provides a Unix-like shell
- Provides a Java API for programming

# Further Readings

- HDFS Architecture
  - <https://hadoop.apache.org/docs/r3.2.2/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>
- Shell commands
  - <https://hadoop.apache.org/docs/r3.2.2/hadoop-project-dist/hadoop-common/FileSystemShell.html>
- FileSystem API
  - <https://hadoop.apache.org/docs/r3.2.2/api/org/apache/hadoop/fs/FileSystem.html>