

# Big Data Processing

# Learning Outcomes

- Get familiar with big-data programming models and query execution engines
- Understand how big-data systems execute user queries internally
- Know the architectural differences between big-data systems

# Big Data Framework

- A system that allows developers to write a program and execute it on a cluster of machines.
- Hides most of the low-level system issues such as fault tolerance, network communication, and load balancing.
- Imposes some restrictions on the developer to ensure that they can run the program efficiently.


# Word Count Example

## Input.txt


```
I often repeat repeat myself
I often repeat repeat
I don't don't know why know why
I simply know that I I I
am am inclined to say to say
a lot a lot this way this way
I often repeat repeat myself
I often repeat repeat
```

Word	Count
a	2
am	2
don't	2
I	9
inclined	1
know	3
lot	2
myself	2
often	4
repeat	8
say	2
simply	1
that	1
this	2
to	2
way	2
why	2

# Word Count Walkthrough (1/2)



# Word Count Walkthrough (2/2)



# Word Count Logic

- The logic behind the word count example can be expressed using only two functions
  - WordExtractor:  $\text{String} \rightarrow \{(w, 1)\}$
  - WordSum:  $(w, \{c\}) \rightarrow (w, \sum c)$

# Complete Word Count in Hadoop

```
public static class TokenizerMapper  
    extends Mapper<Object, Text, Text, IntWritable>{  
private final static IntWritable one = new IntWritable(1);  
private Text word = new Text();  
public void map(Object key, Text value, Context context  
) {  
    StringTokenizer itr = new StringTokenizer(value.toString());  
    while (itr.hasMoreTokens()) {  
        word.set(itr.nextToken());  
        context.write(word, one);  
    }  
}  
}
```

```
public static class IntSumReducer  
    extends Reducer<Text,IntWritable,Text,IntWritable> {  
private IntWritable result = new IntWritable();  
public void reduce(Text key, Iterable<IntWritable> values,  
                  Context context  
) {  
    int sum = 0;  
    for (IntWritable val : values) {  
        sum += val.get();  
    }  
    result.set(sum);  
    context.write(key, result);  
}
```

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

Source: [https://hadoop.apache.org/docs/r3.2.2/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html#Example:\\_WordCount\\_v1.0](https://hadoop.apache.org/docs/r3.2.2/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html#Example:_WordCount_v1.0)

# Complete Word Count in Spark


```
// In Scala shell
val lines = sc.textFile("data.txt")
val pairs = lines.flatMap(s => s.split("\\\\b"))
.map(w => (w,1))
val counts = pairs.reduceByKey((a, b) => a + b)
counts.saveAsTextFile("word_count_output.txt")
```

# Big-data Processing

- Data Model
  - How data is seen by system
- Programming Model
  - How a developer writes a big-data program
- Application Model (Logical Model)
  - How the big-data platform internally represents a user program
- Execution Model (Physical Model)
  - How the program gets executed on the cluster

# Data Model

- Big data systems deal with multiple datasets
- Each dataset is a multiset of records (repeated values)
- Each record can be of any value
- All records in a dataset must have the same type
- Special handling is given when the object is a key-value pair, as detailed later
- A dataset could be partitioned across machines




# Functional Programming Model

- A big-data program consists of user-defined functions
  - E.g., Map and Reduce functions in a Hadoop MapReduce program
- A valid function must satisfy two constraints
  - Stateless/Memoryless
  - Deterministic

# Functional Programming

- Which of these are functions?



# Word Count Functions

- Word Extractor(Line: String) {  
    words = Line.split  
    foreach (w ∈ words) output.write(w, 1)  
}
- SumWords(word: String, counts: int[]) {  
    sum = sum(counts)  
    output.write(word, sum)  
}

# Examples

```
Function1(x) {  
    return x + 5;  
}
```


```
Int sum  
Function2(x) {  
    sum += x;  
    return sum;  
}
```

```
RNG random;  
Function3(x) {  
    random.randomInt(0, x);  
}
```

```
const Map<String, Int> lookuptable;  
Function4(x) {  
    return lookuptable.get(x);  
}
```


# Directed Acyclic Graph

- The functional programming paradigm allows the developer to define one function
- The program consists of multiple functions



# DAG for Word Count in Spark


```
// In Scala shell
val lines = sc.textFile("data.txt")
val wordPairs = lines.flatMap(s => s.split("\\b"))
.map(w => (w, 1))
val wordCounts = pairs.reduceByKey((a, b) => a + b)
wordCounts.saveAsTextFile("word_count_output.txt")
```




# Bulk Synchronous Parallel (BSP)

- The BSP model is how big-data frameworks execute a program
- The model splits the execution into stages of local processing
- Computation stages are separated by a communication barrier


# BSP Model



# Communication Patterns





Fully Connected  
(Requires network communication)




One-to-one  
(Can be done locally in one stage)

# Word Count Stages




# Hadoop MapReduce




# Map and Reduce Functions

- Map Function
  - Maps a single input record to a set (possibly empty) of intermediate records
  - Map:  $\langle k_1, v_1 \rangle \rightarrow \{\langle k_2, v_2 \rangle\}$
- Combine Function (Optional)
  - Combines *some* intermediate records with the same key to a set (possibly empty) of intermediate records
  - Combine:  $\langle k_2, \{v_2\} \rangle \rightarrow \{\langle k_2, v_2 \rangle\}$
- Reduce Function
  - Reduces *all* intermediate records with the same key to a set (possibly empty) of output records
  - Reduce:  $\langle k_2, \{v_2\} \rangle \rightarrow \{\langle k_3, v_3 \rangle\}$


# MapReduce DAG



# MapReduce Program Cycle



# Job Execution Overview



Job  
submission

Job  
preparation

Map

Shuffle

Reduce


Cleanup

# Job Submission


- Execution location: Driver node
- A driver machine should have the following
  - Compatible Hadoop binaries
  - Cluster configuration files
  - Network access to the master node
- Collects job information from the user
  - Input and output paths
  - Map, reduce, and any other functions
  - Any additional user configuration
- Packages all this in a Hadoop Configuration

# Hadoop Configuration

Key: String	Value: String
Input	hdfs://user/eldawy/README.txt
Output	hdfs://user/eldawy/wordcount
Mapper	edu.ucr.cs.bigdata.eldawy.WordCount
Reducer	...
JAR File	...
User-defined	User-defined



Serialized over network




Master node

# Job Preparation

- Runs on the master node
- Gets the job ready for parallel execution
- Collects the JAR file that contains the user-defined functions, e.g., Map and Reduce
- Writes the JAR and configuration to HDFS to be accessible by the executors
- Looks at the input file(s) to decide how many map tasks are needed
- Makes some sanity checks
- Finally, it pushes the BRB (Big Red Button)



# Job Preparation



# Map Phase


- Runs in parallel on worker nodes
- $M$  Mappers:
  - Read the input
  - Apply the map function
  - Apply the combine function (if configured)
  - Store the map output
- There is no guaranteed ordering for processing the input splits

# Input record reader


- Split the file based on the file metadata
  - File size, block sizes, # of nodes
- Each split is defined by:
  - File name, Start offset, Length
- For each split:
  - Seek to the start offset
  - Skip the first record (except for the first split)
  - Read until the beginning of the record goes beyond the start + length

# Map Phase

## Master node



# Input Splits (Map tasks)



# Map Task

- Reads the job configuration and task information (mainly, InputSplit)
- Instantiates an object of the Mapper class
- Instantiates a record reader for the assigned input split
- Reads records one-by-one from the record reader and passes them to the user-defined map function
- Passes the map output to the next step


# Map output

- What happens to the map output?
- It depends on the number of reducers
  - 0 reducers: Map output is written directly to HDFS as the final answer
  - 1+ reducers: Map output is passed to the shuffle phase


# Shuffle Phase

- Executed only in the case of one or more reducers
- Transfers data between the mappers and reducers
- Groups records by their keys to ensure local processing in the reduce phase


# Shuffle Phase



# Shuffle Phase (Map-side)




# Shuffle Phase (Reduce-side)



# Reduce Phase

- Apply the reduce function to each group of similar keys



# Output Writing

- Materializes the final output to disk
- All results are from one process (mapper/reducer) are stored in a subdirectory
- An OutputFormat is used to:
  - Create files in the output directory
  - Write the output records one-by-one to the output
  - Merge the results from all the tasks (if needed)
- While the output writing runs in parallel, the final commit step runs on a single machine.
- The commit step typically writes an empty “\_SUCCESS” file to indicate job success.

# Hadoop MapReduce Conclusion

- A MapReduce program consists of a map, a reduce, and optionally a combine function
- Hadoop distributes the program to all executor nodes
- The input is partitioned and each input split is processed independently
- The intermediate data is shuffled and reduced to produce the final output.