

Intro to Deep Learning

Lecture 5: Backprop and Optimization

Midterm

- Time: Feb. 9 Monday 9:30 - 10:50 am (the usual lecture time)
- Location: Materials Sci and Engineering Room 103 (the usual lecture location)
- Format: 30 multi-choice questions (potentially have multiple correct answers)
- One A4 cheat sheet allowed (hand written!)

- HW1 due this Friday
- HW2 released

Midterm

- Question format:

Problem 1. Which of the following statements regarding course EE/CS 228 are correct?

(A) Lecture time is Tu/Th 8:00 - 9:20 am



MW 9:30 - 10:50

(B) Lecture location is Sproul Hall Room 1340



Olmsted Room 1212

(C) Midterm is on Nov. 9th Thursday (in class)



Feb 9 in class

(D) EE/CS 228 is a great course



Select all correct answers

Midterm covers everything lectured up to Feb 9th

Recall Stochastic Gradient Descent

1. Pick example (x, y)
2. Pick ML model $f_w(\cdot)$ with weights w
3. Pick loss function ℓ to minimize (e.g. squared)
4. Optimization becomes:

$$\mathcal{L}(w) = \ell(y, f_w(x))$$

5. Gradient via chain rule

$$\nabla \mathcal{L}(w) = \ell'(y, f_w(x)) \times \nabla f_w(x)$$

Derivative of Loss


Gradient of Model

HOW?


Loss Functions

- Square loss: $\ell(y, f(x)) = \frac{1}{2}(f(x) - y)^2$
 $\ell'(y, f(x)) = (f(x) - y)$
 - Cross entropy loss
- $$\ell(y, f(x)) = - \sum_{i \in [k]} y_i \log(p_i), \quad \text{where} \quad p_i = \frac{e^{f_i(x)}}{\sum_{j=1}^n e^{f_j(x)}}$$


How to differentiate a neural net?



Answer: *Backpropagate*




Function: $f = x_1 - x_2$




Function: $f = x_1 - x_2$

Gradient: $\frac{\partial f}{\partial x_1} = 1, \frac{\partial f}{\partial x_2} = -1$




Function: $f = x_1 - x_2 = x_1 - (x_3 + x_4)$




Function: $f = x_1 - x_2 = x_1 - (x_3 + x_4)$

Gradient: $\frac{\partial x_2}{\partial x_3} = 1, \frac{\partial x_2}{\partial x_4} = 1$. What about $\frac{\partial f}{\partial x_3}$?




Function: $f = x_1 - x_2 = x_1 - (x_3 + x_4)$

Gradient: $\frac{\partial f}{\partial x_3} = \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial x_3} = -1$




Function: $f = x_1 - x_2 = x_1 - (x_3 + x_5 + x_4)$

Gradient: $\frac{\partial x_2}{\partial x_5} = 1$




Function: $f = x_1 - x_2 = x_1 - (x_3 + x_5 + x_4)$

Gradient: $\frac{\partial f}{\partial x_5} = \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial x_5} = 1$




Function: $f = x_1 - x_2 = x_1 - (w_3 x_3 + w_4 x_4)$




Function: $f = x_1 - x_2 = x_1 - (w_3 x_3 + w_4 x_4)$

Gradient: $\frac{\partial f}{\partial w_3} = \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial w_3} = -1 \times x_3 = -x_3$




Function: $f = x_1 - x_2 = x_1 - \sigma(w_3x_3 + w_4x_4)$




Function: $f = x_1 - x_2 = x_1 - \sigma(w_3x_3 + w_4x_4)$

Let $net_2 = w_3x_3 + w_4x_4$




Function: $f = x_1 - x_2 = x_1 - \sigma(w_3x_3 + w_4x_4)$

Gradient: $\frac{\partial f}{\partial w_3} = \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial net_2} \frac{\partial net_2}{\partial w_3} = -1 \times \sigma' \times x_3 = -\sigma' x_3$



Function $f = x_1 - x_2 = (x_3 + x_5) - \sigma(w_3 x_3 + w_4 x_4)$




$$\text{Function: } f = x_1 - x_2 = (x_3 + x_5) - \sigma(w_3 x_3 + w_4 x_4)$$

$$\text{Gradient: } \frac{\partial f}{\partial x_3} = \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial net_2} \frac{\partial net_2}{\partial x_3} + \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial x_3} = -1 \times \sigma' \times w_3 + 1 \times 1 = -\sigma' w_3 + 1$$

Summary

- Forward to compute f
- Backward to compute the gradients



Activation Functions

- ReLU $\text{ReLU}(x) = \max\{x, 0\}$

$$\text{ReLU}'(x) = \begin{cases} 0, & x \leq 0 \\ 1, & x > 0 \end{cases}$$

Sub-gradient
used at $x = 0$

- Sigmoid $\sigma(x) = \frac{1}{1 + e^{-x}}$


$$\sigma'(x) = \sigma(x) \cdot (1 - \sigma(x))$$

- $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

One-hidden layer network

- $f_\theta(x) = w^\top \sigma(Ux)$ where $\theta = (w, U)$

- $f_\theta(x) = \sum_{i=1}^k w_i \sigma(u_i^\top x)$



One-hidden layer network

- $f_\theta(x) = \sum_{i=1}^k w_i \sigma(u_i^\top x)$
- $\nabla f_\theta(x) = \left(\frac{\partial f}{\partial w}, \frac{\partial f}{\partial U} \right)$
- $\frac{\partial f}{\partial w} = h = \sigma(Ux)$
- $\frac{\partial f}{\partial U} = ?$

One-hidden layer network

- $\frac{\partial f}{\partial u_i} = \frac{\partial w_i \sigma(u_i^\top x)}{\partial u_i} = w_i \sigma'(u_i^\top x) x$

Overall size
 $k \times d$

$$\frac{\partial f}{\partial U} = (w \odot \sigma'(Ux)) x^\top$$

$k \times 1$ vector with
entries $w_i \sigma'(u_i^\top x)$

$1 \times d$ input features

- \odot is Hadamard product (entry-wise multiplication²⁶)

One-hidden layer network

- $\frac{\partial f}{\partial w} = \sigma(Ux)$
- $\frac{\partial f}{\partial U} = (w \odot \sigma'(Ux))x^\top$
- Consider $\mathcal{L}(\theta) = \frac{1}{2} (w^\top \sigma(Ux) - y)^2$
 $\frac{\partial \mathcal{L}}{\partial w} = (w^\top \sigma(Ux) - y) \sigma(Ux)$
 $\frac{\partial \mathcal{L}}{\partial w} = (w^\top \sigma(Ux) - y) (w \odot \sigma'(Ux))x^\top$

Training a one-hidden layer net

Initialize w_0, U_0

Learning rate $\eta > 0$, duration END

for $0 \leq t \leq END - 1$

1. Randomly pick example (x, y) from training data
2. Update weights

- $w_{t+1} = w_t - \eta (w^\top \sigma(Ux) - y) \sigma(Ux)$

- $U_{t+1} = U_t - \eta (w^\top \sigma(Ux) - y) (w \odot \sigma'(Ux)) x^\top$

Return final model $\theta_{FINAL} = (w_{END}, U_{END})$


For HW2, you will train your own neural network from scratch.

Break

Computation Graph

- Optimize

$$J = \ell(W^{(2)}\phi(W^{(1)}x), y) + \|W^{(1)}\|_F^2 + \|W^{(2)}\|_F^2$$



```
# Compute prediction and loss
pred = model(X)
loss = loss_fn(pred, y)

# Backpropagation
loss.backward()
optimizer.step()
optimizer.zero_grad()
```

Computation Graph


- Forward pass:
 - Generate dynamic graph (at least for PyTorch)
 - Store a, z lists
- Backward pass:
 - `loss.backward()` calculate gradients
 - Release memory
- Inference: `torch.no_grad()`
 - No graph will be generated (and no need to store intermediate values)

Gradient Descent Revisit

- Gradient descent iterations

$$f_{i+1} = f_i - \eta \nabla \mathcal{L}(f_i)$$

↑ ↑ ↗ ↗
New model Current model Learning rate Gradient



Linear regression

1. Data: $(x_i, y_i)_{i=1}^n$, $x_i \in \mathbb{R}^d$, $y_i \in \mathbb{R}$
2. Model: Linear model $f(x; w) = x^\top w$
3. Squared Loss: $\ell(f(x), y) = \frac{1}{2} (f(x) - y)^2$
4. Optimize

$$\widehat{w} = \arg \min_{w \in \mathbb{R}^d} \widehat{\mathcal{L}}(w) \quad \text{where} \quad \widehat{\mathcal{L}}(w) = \frac{1}{2n} \sum_{i=1}^n (x_i^\top w - y_i)^2$$

Linear regression

- Let us move to matrix notation

$$X = \begin{matrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{matrix} \quad n \times d \text{ matrix} \qquad y = \begin{matrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_n \end{matrix} \quad n \times 1 \text{ vector}$$

- The loss function becomes

$$\widehat{\mathcal{L}}(w) = \frac{1}{2n} \sum_{i=1}^n (x_i^\top w - y_i)^2 = \frac{1}{2n} \|y - Xw\|_2^2$$

GD for Linear regression

Optimization becomes

$$w^* = \arg \min_{w \in \mathbb{R}^d} \frac{1}{2} \|y - Xw\|_2^2$$

- Calculate the gradient

$$\nabla \mathcal{L}(w) = X^\top X w - X^\top y$$

- Gradient iterations are given by

$$w_{t+1} = w_t - \eta \nabla \mathcal{L}(w_t)$$

GD for Linear regression

Plugging in the gradient formula

$$w_{t+1} = w_t - \eta(X^\top X w_t - X^\top y)$$

Questions:


- Which learning rate to use?
- How many iterations needed?
- Can we speed up further?

Answers are important for **deep learning** as well.

Convex Optimization

- A (differentiable) function $f: \mathbb{R}^d \rightarrow \mathbb{R}$ is convex iff, for any $x_1, x_2 \in \mathbb{R}^d$, we have

$$f(x_1) \geq f(x_2) + \nabla f(x_2)^T (x_1 - x_2)$$



Convex Optimization

- A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is L -smooth if for any $x_1, x_2 \in \mathbb{R}^d$

$$\|\nabla f(x_1) - \nabla f(x_2)\|_2 \leq L\|x_1 - x_2\|_2$$

- L -smooth implies that for any $x_1, x_2 \in \mathbb{R}^d$

$$f(x_1) \leq f(x_2) + \nabla f(x_2)^\top (x_1 - x_2) + \frac{L}{2} \|x_1 - x_2\|_2^2$$

GD for Convex Functions

- Let $f: \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex function that is also L -smooth
- Let $w^\star := \arg \min_w f(w)$ — **the minimizer**
we try to find
- Consider the following gradient descent update

$$w_{t+1} = w_t - \eta \nabla f(w_t)$$

Step I: The Descend

- If $\eta \leq 1/L$, we then have

$$f(w_{t+1}) \leq f(w_t) - \frac{\eta}{2} \|\nabla f(w_t)\|_2^2$$

Function value decreases
strictly unless $\nabla f(w_t) = 0$

Step II: Convergence Rate

We have

- $f(w_{t+1}) - f(w^*) \leq \frac{1}{2\eta} (\|w_t - w^*\|_2^2 - \|w_{t+1} - w^*\|_2^2)$
- $f(w_T) - f(w^*) \leq \frac{\|w_0 - w^*\|_2^2}{2\eta T}$
- To achieve ε error, we need $O\left(\frac{1}{\eta\varepsilon}\right) = O\left(\frac{L}{\varepsilon}\right)$ number of iterations

GD for Linear regression

Optimization becomes

$$f(w) = \frac{1}{2} \|y - Xw\|_2^2$$

- f is convex (actually quadratic)
- It's gradient equals to

$$\nabla f(w) = X^\top X w - X^\top y$$

GD for Linear regression

- It's gradient equals to

$$\nabla f(w) = X^\top X w - X^\top y$$

- $\nabla f(w)$ is $\lambda_{\max}(X^\top X)$ -Lipschitz

$$|\nabla f(w_1) - \nabla f(w_2)| = |X^\top X(w_1 - w_2)| \leq \lambda_{\max}(X^\top X) \cdot \|w_1 - w_2\|_2$$

Select the learning rate as $\eta = 1/\lambda_{\max}(X^\top X)$ and run for $O(\lambda_{\max}(X^\top X)/\varepsilon)$ iterations ensures convergence up to an ε -optimal minimizer.